

AX

AXON-X · AXON-WATCH

How-To *Handbook*

Your field guide for bootstrapping, verifying, troubleshooting,
and shipping thin slices in the next-generation operator console.

Shell :4173

Control plane :8787

Watch :8788

Last verified · 2026-07-29 — IDE Soft Attention actions (**Try
again** / **Explain** / **Open team**) live on the agent review strip
with **Review N files**, **Try again** hides after a successful shift.
Claude Code CLI is a local runtime target. After every push:
`./scripts/ops/watch-fast-gate.sh`.

Generated 10 Aug 2026

CONSOLE-WEB

CONTROL-PLANE

SHARED-TYPES

56 chapters

1

Table of Contents

2

Handbook map

3

Quick Start

4

AXON-X Constitution registries

5

Runtime auth, CLI, and tools

6

VAXON Desktop

7

Terminology And Abbreviations

8

What Axon-X Is

9

What State The Repo Is In Right Now

10

Source Of Truth Rules

11

Repo Layout

12

The Current Architecture In Plain English

13

Detailed setup (first install)

14

What The Current App Does On Boot

15

Locked Shell Layout

16

Current Shell Layout

17

Mockup / Live Parity Notes

18

The Most Important Files Right Now

19

Verification Commands

20

TEST-0 acceptance (``workspace_smoke``)

21

Latency evidence (D1)

22

What PASS / PENDING / FAIL Mean In Verification

23

Common Working Patterns

24

Troubleshooting

25

Debugging playbook

26

Tips And Tricks

27

Snippet cookbook

28

Upgrading and updating

29

What A Good Next Slice Looks Like

30

Final Guidance

31

CI, merge workflow, and worker agents

32

or follow a known run id:

33

Typical Fast Gate early fail: file-size ratchets in scripts/guardrails/hotspot_budgets.json

34

CI locally

35

Open PR

36

PR checks

37

DashPro self-hosted runner (this PC)

38

Worker scheduler tests

39

Stale reconcile (when loaded on CP)

40

Gate 9 — CI remediation (unaware-operator loop)

41

Autonomy gates & service identity (operator directives)

42

~/config/axon-watch/deployment.env

43

unlock Vault in the UI (Remember me), then:

44

~/config/axon-watch/deployment.env

45

/vault → Enable auto-unlock (keyfile already present will auto-unlock on next start)

46

Cloudflare API token with Account → Cloudflare Tunnel → Edit

47

writes `~/config/axon-watch/mtls/{ca,client,server}.{crt,key}`

48

then push and confirm Axon-X Fast Gate is green on GitHub

49

Recent operator features (Gate 3–5 + console UX)

50

From repo root — polls the latest Fast Gate run for this branch

51

Auto-loop status & credits budget (operator brief)

52

Company hierarchy, Lead check-in, and monitoring close-out

53

Recovering a stuck agent and running School Operations Phase 1

54

Reliability and deliberate controls

55

Online research: SearXNG and legacy Google

56

then set
`AXON_WATCH_SEARXNG_URL=http://127.0.0.1:8080`
in `.env` (or vault)

Read this first

Practical guide for operators, reviewers, and developers working in the **axon-watch** repo — the greenfield home of **Axon-X**.

- Plans live in `axon-local/Plans/Axon-Watch/` — implement here, don't redefine semantics there.
- Start with `./scripts/dev/up.sh`, verify with `npm run verify`.
- Treat bootstrap-real and feature-real as different things — see parity ledger before assuming parity.

01 Table of Contents

1. [Quick Start](#) — first 5 minutes
2. [Handbook map](#) — who reads what
3. [Operator manual](#) — daily rituals 3.5. [Runtime auth, CLI, and tools](#) — Pro vs API key, native vs Cursor 3.6. [CI, merge, and worker agents](#) — Fast Gate, `dev`, roster; DashPro self-hosted CI + RAM; canary OTA vs EAS Build; [Gate 9 CI remediation](#) 3.65. [Autonomy gates & service identity](#) — Gate 4 tasks, scheduler off, watch token + mTLS 3.66. [Recent operator features](#) — task board, concurrent tabs, galaxy labels, Lead planner, CI watch 3.67. [Auto-loop status & credits](#) — are we autonomous yet? Cursor / API budget for multi-project 3.68. [Company hierarchy & Lead check-in](#) — VAXON attend loop + AUTONOMOUS ON safety matrix 3.69. [Stuck agent recovery & School Operations Phase 1](#) — restart decisions, VAXON's recovery contract, and daily homework approval 3.70. [AXON-X Constitution registries](#) — missions, evidence, decisions, capabilities, ADRs, debt, health, and self-healing proof 3.7. [VAXON Desktop](#) — packaged Linux install
4. [Teaching Axon-X](#) — explain it to others
5. [Codebase in plain English](#) — what happens under the hood
6. [Source index](#) — where truth lives
7. [Snippet cookbook](#) — copy-paste commands
8. [Terminology](#)
9. [Architecture & repo layout](#) — structure and ownership
10. [Detailed setup](#) — first install
11. [Boot flow](#) — what loads on refresh
12. [Shell layout](#) — regions and modes
13. [Key files](#) — start reading here
14. [Verification](#) — gates and PASS/PENDING/FAIL
15. [Working patterns](#) — how to add code safely

16. [Debugging playbook](#) — step-by-step fixes
17. [Tips, hints & tricks](#)
18. [Upgrading & updating](#) — pulls, deps, planning sync
19. [Next slices](#) — what to build next

02 Handbook map

AUDIENCE	START HERE	THEN READ
Operator (daily use)	Quick Start	Runtime auth, CLI, and tools , VAXON Desktop , Operator manual
Teacher / reviewer	Teaching Axon-X	Verification , <code>docs/FINAL_PARITY_VERIFICATION.md</code>
Developer	Codebase in plain English	Source index , Common working patterns
Integrator / merge	CI, merge, and worker agents	<code>docs/CI_GATES.md</code> , <code>./scripts/ops/watch-fast-gate.sh</code>
Autonomy / remote host	Auto-loop status & credits	Autonomy gates & service identity , Recent operator features
Stuck agent / school operator	Stuck agent recovery & School Operations Phase 1	Debugging playbook , Company hierarchy & Lead check-in
Lead / autonomous agent	AXON-X Constitution registries	Verification , <code>docs/ops/agent-reports/axon-x-constitution-gap-audit-2026-08-09.md</code>
Debugger	Debugging playbook	Troubleshooting
Upgrader	Upgrading & updating	<code>./scripts/ops/sync_planning_mirror_to_axon_local.py</code>

Shorter onboarding: `docs/AXON-X-STARTER-GUIDE.md`

03 Quick Start

This section is the living operator onboarding guide.

Start the stack

Always-on host (this machine — preferred):

COMMAND	WHAT IT DOES
<code>axonhealth</code>	Probe console + control-plane + watch (+ key APIs)
<code>axonrestart</code>	Soft restart of systemd user units, then health check
<code>axonrevive</code>	Use when the shell is empty (Runtime unavailable / No workspace). Force-kills a wedged control-plane, restarts all three units, health-checks
<code>axonfixconnectors</code>	Use when Mission Control shows REQUIRED CONNECTOR DOWN / tunnel token missing / vault unlock HTTP 503 about <code>AXON_WATCH_INTERNAL_SERVICE_TOKEN</code>

```
axonhealth      # is everything up?
axonrevive      # empty shell / hung API – fix it
axonfixconnectors # required connector / tunnel / vault 503 – diagnose (+ optional fix)
# then hard-refresh http://127.0.0.1:4173
```

These are on your PATH (`~/local/bin` → `bin/` in this repo). See [Snippet cookbook](#).

Reliability and deliberate controls

See [Reliability and deliberate controls](#) for Vite recovery, manual STAND-UP, speech onset, and AgentDock hover actions.

Dev bootstrap (alternate, not used when systemd owns the ports):

```
cd /run/media/vaxon/axon-data/repos/axon-nvme/repos/axon-watch
./scripts/dev/up.sh
./scripts/dev/check-health.sh
```

Open **http://127.0.0.1:4173** and hard-refresh after upgrades (`Ctrl+Shift+R`).

Important: `./scripts/dev/down.sh` does **not** stop systemd always-on units. On this host use `axonrestart` / `axonrevive`.

04 AXON-X Constitution registries

The Constitution work adds a durable executive memory layer for AXON-X. It does **not** replace the original execution stores; it indexes them so VAXON, Dana, specialists, and the operator can answer: what mission is active, what evidence exists, what decision was made, what capability owns it, and what platform health looked like at the time.

What changed

The control plane now has these SQLite-backed registries:

REGISTRY	WHAT IT MEANS	WHY IT MAKES AXON-X STRONGER
Evidence	Pointers to run history, Lead receipts, autonomy receipts, deliveries, and host-action receipts	Agents stop inventing receipts and can point to source records.
Missions	Operator goals with workspace, risk, Lead-plan link, checkpoint, and success criteria	Lead work becomes durable mission progress, not only chat text.
Decisions	Autonomous/operator-gated decisions with actor, tier, risk, confidence note, and evidence IDs	“Why did VAXON/Dana do that?” has a queryable answer.
Capabilities	Stable <code>CAP-###</code> anchors for major platform abilities	Blockers can name the exact capability that is missing or degraded.
ADRs	Canonical <code>docs/adr/ADR-*.md</code> records	Architecture choices are visible to agents without rereading the repo.
Technical debt	Known gaps with severity, area, and evidence links	Gaps survive handoffs instead of being rediscovered.
Platform health	Runtime/watch/degraded snapshots	Self-healing decisions can prove current service posture first.

How agents should use it

- Before claiming a dispatch, recovery, or blocker is real, look for evidence through `/api/operator/constitution/evidence`.

- When a Lead plan is persisted, the system creates a linked mission automatically and attaches Lead receipt evidence to that mission.
- When VAXON/autonomous attention records a safe action, the system records both evidence and a decision row.
- If platform health matters, capture it with `POST /api/operator/constitution/health/capture-runtime-summary`.
- If seed data is missing, run `POST /api/operator/constitution/seed`; this indexes canonical ADR markdown and seeds stable capability anchors such as `CAP-007`, `CAP-034`, and `CAP-070`.

Operator console surface

Open **Settings** → **Constitution** to view the registry counts and recent human-readable summaries for missions, decisions, capabilities, ADRs, debt, and health. This surface is read-only on purpose: it gives visibility without letting a UI refresh mutate planning state.

Verification

Use:

```
npm run verify:constitution
python -m pytest tests/test_constitution_registry.py tests/test_constitution_gate_script.py
-q
npm run test -w @axon-watch/console-web -- constitution-overview-view
```

`verify:constitution` checks the registry spine, required routes, ADR/capability seed anchors, auth coverage for mutating routes, focused tests, and the handoff ledger.

Pick a real workspace

The left sidebar should show **axon-watch** and **axon-local** (bound project roots). Start with **axon-watch** for Axon-X development, or **axon-local** when you need the legacy repo context.

Demo names like `workspace_nlp` are mock catalog entries — they are hidden when real project bindings exist.

Two modes — different jobs

Axon-X has two layout modes (top-right toggle). They are **not** two different apps; they are two views over the same workspace, runs, and APIs.

	OPERATOR MODE	IDE MODE
Purpose	Run oversight, signals, approvals, command execution	Files, editor, terminal, agent dock
Center	Mission control — run phase, live feed, resume/complete	Monaco editor + bottom terminal dock
Left sidebar	Workspaces + Attention (signals, inbox, receipts)	Explorer / search / git activity bar
Right dock	Conversation transcript + Command/KAIRO hero	Resizable agent dock (conversation + composer)
Best for	“What is running? What needs me? Run this command.”	“Edit files, use terminal, review code.”
Input style	Exact commands in the Command seam (see footer Commands)	Same command seam in agent dock + full editor/terminal

Operator mode is the default production surface for day-to-day oversight.

IDE mode is for hands-on work in the bound repo (real `project_root` on disk).

Switch modes anytime — workspace selection, runs, and conversation thread persist.

What you can do in Axon-X today (v1)

Real and verified today:

- Select **axon-watch** or **axon-local** workspace
- View **runtime summary**, **inbox signals**, and **KAIRO briefing** from live APIs
- Track **run phase** in mission control (stop/resume/review-ready flows)
- Send **supported commands** (not free-form chat) via the Command seam
- Run **git status** against the bound repo root
- **IDE mode**: open workspace files in Monaco, PTY terminal in repo root
- **Attention sidebar**: connector/signal/delivery visibility

Still thin or deferred (use axon-local `:7734` fallback if needed):

- General conversational chat (“Hi”, “explain this repo”)
- Full agent tool loop parity with classic Axon
- Child-project connectors and legacy integration surfaces
- Native tray notifications beyond hide-on-close packaging

Supported commands (Operator mode)

The conversation/command seam accepts **exact commands** only. Natural language will return “unsupported command”.

COMMAND	WHAT IT DOES
<code>health</code> / <code>api/health</code>	Probe control-plane health
<code>ls</code> / <code>list files</code>	List files in the workspace
<code>read README.md</code> / <code>cat notes.txt</code>	Read a workspace file
<code>git status</code>	Git status in the bound project root
<code>resume from review</code>	Resume the primary <code>review_ready</code> run

In the UI: footer **Commands** button (Operator mode) opens the list and can prefill the Command seam. Source of truth: `apps/console-web/src/lib/operator-supported-commands.ts` (keep in sync with `services/control-plane/app/chat/command_executor.py`).

Typical first session

1. Open `:4173` → wait for boot overlay → shell loads
2. Left sidebar → select **axon-watch**
3. Operator mode → right dock → **Command** tab
4. Type `git status` → send → read agent receipt in **Conversation**
5. Center mission control shows run phase if a run was created
6. Toggle **IDE** → open `README.md` , use terminal in the real repo root
7. Left **Attention** → inspect signals when watch reports degraded summary

When things look noisy

Dev SQLite may contain old smoke runs (“32 runs ready for review”). Reset if needed:

```
./scripts/dev/down.sh
rm -f .local/state/control-plane.sqlite3
./scripts/dev/up.sh
```

Verify smoke

```
npm run verify:production-operator
```

Understanding runs, review_ready, RESUME, and COMPLETE

Analogy: Axon-X is a **supervised assistant**, not a chat bot. When you send a command like `read README.md` or `health`, the system starts a **run** (a tracked job): do the work → show output in **Conversation** → **pause** at a checkpoint called **review_ready**. That pause is intentional — you look at the result before anything else happens.

BUTTON / ACTION	PLAIN MEANING	WHEN TO USE IT
COMPLETE RUN	“I’m done — this job succeeded.”	One-shot commands (<code>read ...</code> , <code>git status</code> , <code>health</code>) when output looks good
RESUME RUN	“I’ve reviewed it — keep this run going.”	Multi-step work, or when KAIRO ADVISE says to resume a named run
Command → <code>resume from review</code>	Same as RESUME for the primary paused run	When buttons are unclear; acts on one run at a time
Left → Attention	Signal inbox (bootstrap noise, delivery receipts)	When footer shows SIGNALS: N ACTIVE — usually informational in dev

“Phase is now review_ready. Review when ready.” — not an error. The command finished; Mission Control is asking you to **COMPLETE** (usual) or **RESUME** (if more steps expected).

Mission Control AUTONOMOUS (bounded)

In **OPERATOR** → **Mission Control**, use **AUTONOMOUS ON/OFF** for bounded worker control. See [Auto-loop status & credits](#) for the complete safety matrix, approval behavior, dedupe rules, and emergency procedure.

“2 runs are ready for operator review” — what that means

This is a **count of paused jobs**, not a failure. Each command you ran (`health` , `read README.md` , `git status` , ...) can leave its own run sitting in **review_ready** until you **COMPLETE** or **RESUME** it. **2 runs** = two unfinished checkpoints (e.g. one Health run + one Read README run).

What to do:

1. **Center** → **Mission Control** — handles the **primary** (most recent) run via **RESUME RUN** / **COMPLETE RUN**.
2. **Right** → **KAIRO Briefing** — see **NOTICE** (the count) and **ADVISE** (suggested next click).
3. **Clear the backlog** — for each run you’re happy with, click **COMPLETE RUN**. Repeat until the notice says “No active runs” or only one remains.
4. **Dev reset** (optional) — wipe stale runs from smoke tests:

```
./scripts/dev/down.sh && rm -f .local/state/control-plane.sqlite3 && ./scripts/dev/up.sh
```

KAIRO NOTICE / ADVISE / DECIDE — where to go

KAIRO Briefing (right dock → **KAIRO** tab, or footer **Open KAIRO Briefing**) is a **short summary**, not a second app. It mirrors the same backend truth as Mission Control and Attention.

LABEL	MEANING	WHERE TO ACT
NOTICE	Headline (“2 runs are ready...”)	Read only — context
ADVISE	Suggested next step (e.g. “Resume Health.”)	Usually → RESUME RUN in center, or COMPLETE RUN if that job is done
DECIDE	What choice is waiting	Center buttons, or Attention for signals
EXECUTE	Concrete action phrase	Command tab, or the button ADVISE points to

Example: ADVISE says “**Resume Health.**” → you previously ran `health` and that run is paused → go to **Mission Control** → **RESUME RUN** (continues that run) or **COMPLETE RUN** (closes it if you only wanted a one-time health check).

Signals (e.g. “**Bootstrap: runtime summary stale**”) in **Attention** are often **expected in local dev** — watch/bootstrap scaffolding, not production outage. Review them in **Left** → **Attention**; they do not block **COMPLETE RUN** unless an approval gate is open.

Bootstrap & signals — what to do

Bootstrap in Axon-X means the console and services are up, but some runtime summary fields are still intentionally thin while watch/control-plane parity grows. That is normal in local dev — not the same as “the app is broken.”

WHAT YOU SEE	WHAT IT MEANS	WHAT TO DO
Bootstrap: runtime summary stale (Attention → Signals)	Watch is connected; summary assembly is bootstrap-thin	Ignore , tap Details , or hit CLEAR to acknowledge and hide bootstrap noise
OBSERVE chip on a signal	KAIRO watch mode; informational only	No click action. Read-only label.
DELIVERED chip	Delivery receipt was recorded	No click action. Read-only label.
SIGNALS: N ACTIVE (footer)	N open inbox signals	Open Attention or Open Attention from Mission Control — review, then return to Command
IDLE + bootstrap signal only	Nothing waiting on you	No run action required — optional signal review only

When bootstrap noise is OK: local `./scripts/dev/up.sh`, watch healthy, no pending approvals, runs complete normally.

When to investigate: approvals stuck open, runs won't resume/complete, watch disconnected in footer, or bootstrap signal persists **after** a production deploy with full runtime summary expected.

See also [Tip 6: bootstrap-real vs feature-real](#).

Run names — what you see in the UI

Every command creates a **run** with two identifiers:

WHAT YOU SEE	MEANING
Health check, Read README.md, Git status	Friendly task name (from your command)
#cdb931 (6 characters)	Short run ref — internal tracking only; you rarely need the full <code>run_...</code> id

Old runs stored raw command text (`health` , `git status`); the UI **humanizes** those labels. KAIRO **ADVISE** uses the same friendly names (e.g. “Resume Health check.” not “Resume health.”).

When **2+ paused tasks** appear, Mission Control lists each friendly name — click **COMPLETE RUN** for the current one, repeat until the queue clears.

05 Runtime auth, CLI, and tools

This section explains how Axon-X authenticates agent runtimes, when you need secrets in **/vault**, and when Axon uses its own executors vs external CLIs.

Official Cursor CLI reference: [Cursor CLI authentication](#)

Two auth paths for Cursor (Pro / daily use vs headless)

PATH	WHEN TO USE	WHAT YOU DO	VAULT KEY REQUIRED?
CLI subscription (recommended)	Daily operator work on your machine with Cursor Pro/Team	Run <code>cursor agent login</code> once on the host ; verify with <code>cursor agent status</code>	No — <code>CURSOR_API_KEY</code> is optional
API key (headless / CI)	Servers without browser, automation, cloud agents	Generate key in Cursor Dashboard → Integrations → API Keys ; store as <code>CURSOR_API_KEY</code> in /vault or shell env	Yes (or shell env)

Axon-X probes subscription auth live:

- **Vault consumer** (`/vault` → Consumer readiness): marks **Cursor CLI runtime Ready** when `cursor agent status` shows a logged-in account, even with zero vault keys.
- **Control plane** (`GET /api/runtime/status` , `GET /api/runtime/cursor/status`): same probe; composer shows account + auth method.
- **Dispatch**: if subscription is active, control-plane **strips** `CURSOR_API_KEY` from the subprocess env to avoid auth conflicts (browser login and API key are alternate paths per Cursor docs).

Pro without vault key: if `cursor agent status` prints `Logged in as ...@...` and vault search shows no `CURSOR_API_KEY` , you are on the **subscription path** — correct for daily Pro use.

GitHub CLI (`gh`) for draft-PR delivery

Worker shifts that finish with `push_policy=draft_pr` need the **GitHub CLI** on the **control-plane host** (not inside the chat UI). If Team shows `Delivery blocked: gh CLI is required to open a draft PR` , fix the host:

1. Install: <https://cli.github.com/> (or `sudo apt install gh` / package manager).
2. Auth once as the operator user: `gh auth login` (HTTPS or SSH matching the repo remotes).
3. Confirm control-plane can see it: - `command -v gh` and `gh auth status` - If systemd PATH omits `~/.local/bin` , either restart after `scripts/ops/run-service.sh` (it prepends `~/.local/bin`) **or** set `AXON_WATCH_GH_CLI_PATH=/absolute/path/to/gh` in `~/.config/axon-watch/deployment.env` , then `systemctl --user restart control-plane` .
4. Retry the teammate (**Try again** / **Continue** on the agent review strip above the composer — only when the last job failed or was interrupted; succeeded shifts hide that CTA).

Source: `services/control-plane/app/workspace_delivery/gh_cli.py` , `publish.py` (`_open_or_update_draft_pr`).

Codex / OpenAI auth

PATH	SETUP
Codex CLI login	<code>codex login</code> on the host; vault consumer probes <code>codex login status</code>
API keys in vault	<code>CODEX_API_KEY</code> or <code>OPENAI_API_KEY</code> in /vault (either satisfies the codex consumer)

Claude Code auth

PATH	SETUP
Claude Code CLI login	<code>claude auth login</code> on the host; vault consumer probes <code>claude auth status</code>
API key in vault	Optional <code>ANTHROPIC_API_KEY</code> for headless; subscription login preferred for Max/Pro

Select **Claude Code CLI (local)** in the composer runtime picker (or set `AXON_WATCH_IDE_RUNTIME_TARGET=claude_local`) so workers burn Claude usage instead of Cursor credits.

Vault consumers vs runtime dispatch

Consumers are readiness labels for operators — they never expose secret values.

CONSUMER	READY WHEN	OPTIONAL / FALLBACK
<code>cursor_runtime</code>	CLI subscription or <code>CURSOR_API_KEY</code> in vault	API key only needed for headless
<code>claude_runtime</code>	<code>claude auth login</code> or <code>ANTHROPIC_API_KEY</code> in vault	Max/Pro subscription preferred
<code>codex_runtime</code>	<code>codex login</code> or Codex/OpenAI vault keys	
<code>openai_provider</code>	<code>OPENAI_API_KEY</code> in vault	Direct OpenAI fallback
DashPro monitor consumers	Required monitor keys in vault/import	

Source: `services/axon-watch/app/vault/snapshot.py` , `cli_runtime_probe.py`

Runtime dispatch (actually running a model) lives in the control-plane:

- `services/control-plane/app/cli_runtime/catalog.py` — auth probes, ready flags
- `services/control-plane/app/cli_runtime/router.py` — picks binary, env, retry without API key on oauth
- `services/control-plane/app/cli_runtime/vault_keys.py` — merges unlocked vault keys into runtime context

Native Axon tools vs Cursor CLI

Axon-X has **two execution lanes** — do not mix them when debugging.

SURFACE	LANE	EXECUTOR	TOOLS
Operator mode → Command seam	Lane A — Command	Axon <code>command_executor.py</code>	Exact commands only (<code>health</code> , <code>git status</code> , <code>read ...</code>) — Axon native
IDE mode → Agent dock composer	Lane B — Ask / Plan / Agent	Cursor CLI subprocess (<code>cursor</code> <code>agent ...</code>)	Cursor model + consultative prompt; not full Cursor IDE agent tools yet
Future (G3.5+)	MCP registry	Planned wiring	Static registry exists; not connected to dispatch today

When to use native Axon tools: Operator oversight, deterministic repo commands, health checks, file reads — anything in the [supported commands](#) table.

When Cursor CLI runs: IDE composer messages in Ask, Plan, or Agent mode with runtime target **Cursor CLI (local)**. Streaming (SSE) applies to this lane when `AXON_WATCH_LANE_B_STREAMING=1` (default).

Composer modes (IDE dock):

MODE	CURSOR CLI FLAG	BEHAVIOR TODAY
Ask	<code>--mode ask</code>	Read-only style answers
Plan	<code>--mode plan</code>	Step mapping before execution
Agent	Default consultative; Full Access in composer → approval → Cursor <code>--mode agent</code> / Codex workspace-write	<code>execution_access: full</code> + G3.3 approval gate

Agent review strip (above the composer input):

CONTROL	WHEN IT APPEARS
N files / Review N files	Agent edited files in the current thread
Try again / Continue	Soft Attention only — last teammate job failed or was interrupted . Hidden after a successful shift
Explain / Open team	Same Soft Attention window (Explain needs a linked run receipt)
Stop / Resume / Apply all	Live run control and review-ready apply

The soft Attention **Try again** / **Explain** / **Open team** actions live on this strip with **Review N files**. Succeeded shifts hide the retry CTA.

Choosing runtime target and model

Open the **model picker** (⚡ chip) in the IDE agent dock:

1. **Runtime target** — `Cursor CLI (local)`, `Claude Code CLI (local)`, `Codex CLI (local)`, or cloud placeholders. Preference persists in shell local storage via `shell.setSelectedRuntimeTarget`.
2. **Auto toggle** — when ON, Cursor picks the best model per request (no `--model` flag). When OFF, your pinned catalog model is passed to the CLI.
3. **Add models** — browse live output of `cursor agent --list-models` (cached ~5 min). Badges like **Fast** / **High** come from CLI labels when present.
4. **Auth line** — shows `CLI subscription · you@domain` or vault/API-key status; **Open Vault** only when auth is actually blocked (vault locked or missing keys **and** CLI not signed in).

Environment overrides:

VARIABLE	PURPOSE
<code>AXON_WATCH_CURSOR_CLI_PATH</code>	Non-default <code>cursor</code> binary path
<code>AXON_WATCH_CLAUDE_CLI_PATH</code>	Non-default <code>claude</code> binary path
<code>AXON_WATCH_CODEX_CLI_PATH</code>	Non-default <code>codex</code> binary path
<code>AXON_WATCH_IDE_RUNTIME_TARGET</code>	Force default target (<code>claude_local</code> , <code>cursor_local</code> , ...)
<code>AXON_WATCH_IDE_RUNTIME_FAMILY</code>	Prefer family when choosing default (<code>claude</code> , <code>cursor</code> , <code>codex</code>)
<code>AXON_WATCH_LANE_B_STREAMING</code>	<code>1</code> (default) SSE streaming for IDE composer; <code>0</code> in tests

API endpoints:

- `GET /api/runtime/status` — all targets + vault posture
- `GET /api/runtime/cursor/status?force_refresh=1` — Cursor auth + live model catalog
- `GET /api/vault/status` — consumer readiness (axon-watch service)

Troubleshooting auth

SYMPTOM	LIKELY CAUSE	FIX
Vault shows Missing: CURSOR_API_KEY but CLI is logged in	Stale UI or old snapshot	Refresh /vault; consumer should show CLI subscription · ...
Composer says invalid API key	Stale <code>CURSOR_API_KEY</code> in vault or shell env conflicting with subscription	Remove bad key from vault; unset shell <code>CURSOR_API_KEY</code> ; restart control-plane
<code>cursor agent status</code> → not logged in	No subscription session on host	<code>cursor agent login</code>
Runtime ready but dispatch fails	Wrong binary or model id	Check <code>AXON_WATCH_CURSOR_CLI_PATH</code> ; pick Auto or a model from live catalog
Open Vault when Pro works	Should not show if oauth ready	Hard-refresh; verify <code>GET</code> <code>/api/runtime/status</code> shows <code>ready:</code> <code>true</code>
Out of usage / rate limit from Cursor	Current CLI account hit subscription quota	Switch accounts (below) or use Auto / another model; admin may need to raise team limits

Switch Cursor subscription account (logout → login)

When the agent reports **out of usage** or you need a different Pro/Team account on this machine:

```
# 1. Sign out (clears stored CLI session on this host)
cursor agent logout

# 2. Confirm logged out
cursor agent status  # should say not logged in / auth required

# 3. Sign in with the other Cursor account (opens browser)
cursor agent login

# 4. Verify the new account
cursor agent status  # expect: Logged in as other@domain
```

Headless terminal (no browser auto-open):

```
NO_OPEN_BROWSER=1 cursor agent login
# Open the printed URL manually in a browser logged into the target account
```

After switching:

1. Refresh **/vault** → Cursor consumer should show the new **CLI subscription · email**

2. Hard-refresh `:4173` or reopen the model picker (forces runtime status refresh)

3. Retry IDE composer — old thread errors are from the prior account/session; start a new turn

Note: Axon-X does not store Cursor passwords. Logout/login only affects the **host Cursor CLI** session. Vault `CURSOR_API_KEY` (if present) is separate — remove or update it if you switch to API-key auth instead.

Quick host checks:

```
cursor agent status          # expect: Logged in as ...  
echo "${CURSOR_API_KEY:+set}" # empty = good for Pro path  
curl -s http://127.0.0.1:8787/api/runtime/cursor/status | python3 -m json.tool
```

06 VAXON Desktop

See <docs/how-to/vaxon-desktop.md>.

07 Terminology And Abbreviations

Use this glossary when reading plans, ADRs, code, or agent summaries.

TERM	MEANING
ADR	Architecture Decision Record — a numbered, immutable write-up of a significant technical or process choice. Accepted ADRs live in <code>docs/adr/</code> . Do not rewrite them; supersede with a new ADR instead.
Axon-X	User-facing product name for the next-generation operator console.
axon-watch	Internal repo folder and npm workspace name. Legacy naming; not the product label shown to operators.
axon-local	The current production Axon app repo (port 7734). Planning for Axon-X still lives here under <code>Plans/Axon-Watch/</code> .
Bootstrap	Early boot / dev scaffolding state — services run, but some DTO fields (runtime summary depth, signal richness) are intentionally thin until parity slices land. Signals like <code>Bootstrap: runtime summary stale</code> are expected locally , not production outages.
Briefing seam	<code>GET /api/briefing</code> returns canonical <code>OperatorBriefing</code> . The shell loads it at bootstrap and projects that data across the right dock: approvals, signals, and the KAIRO briefing card all read from the same briefing/runtime truth. Approval mutations stay on the run approval seam. See <code>docs/contracts/BRIEFING-SEAM.md</code> .
Control plane	FastAPI service on port 8787 that owns run truth, runtime summary, inbox projection, workspaces list, and briefing.
Console-web / shell	Vue 3 frontend on port 4173 — the visible Axon-X UI.
Contract / shared contract	Canonical TypeScript types and JSON fixtures in <code>packages/shared-types/</code> . Frontend and backend must agree here first.
Coordinator lane	Single owner of serial-only semantics during multi-agent work. See <code>docs/MULTITASK-LANES.md</code> .
DTO	Data Transfer Object — a typed payload shape exchanged between services or UI layers (for example <code>RuntimeSummary</code> , <code>RunRecord</code>).
Fitness function	An automated check that guards architecture or performance (dependency direction, DTO size budgets, latency thresholds).
Frozen planning bundle	The locked docs under <code>axon-local/Plans/Axon-Watch/</code> . Implementation must not silently drift from these.
KAIRO	Knowledge-Augmented Intelligence for Response and Oversight — operator-presence layer (watching, advising, interrupting, executing with receipts). NOTICE / ADVISE / DECIDE in the KAIRO Briefing panel are rhythm labels from <code>GET /api/briefing</code> — suggested reading, not separate commands. See Quick Start → <i>KAIRO NOTICE / ADVISE</i> .

TERM	MEANING
Lane A/B/C/D	Parallel implementation ownership areas defined in <code>docs/MULTITASK-LANES.md</code> (watch, shell, control-plane, dev/verify).
Monaco host	In-browser code editor surface (<code>EditorHost.vue</code>). Loads workspace files on disk (README.md, notes.txt) plus read-only DTO overview tabs.
Parity ledger	Checklist of behaviors Axon-X must eventually match from current Axon. Lives in frozen planning.
Run / run record	Canonical execution unit with <code>phase</code> , <code>status</code> , capability flags, and transition history.
Run phase	High-level lifecycle stage (<code>queued</code> , <code>starting</code> , <code>executing</code> , <code>awaiting_approval</code> , <code>review_ready</code> , <code>completed</code> , etc.). Defined in frozen <code>run-state.md</code> .
Runtime summary	Boot-critical DTO at <code>GET /api/runtime/summary</code> — identity, watch connection, active runs, approvals snapshot.
Signal / inbox item	Watch-produced event surfaced through control-plane <code>GET /api/inbox</code> .
Thin slice	A small, verifiable vertical increment — one owned behavior with tests, not a broad rewrite.
Watch service	FastAPI service on port 8788 that produces canonical signals; control-plane projects them into inbox/runtime summary.
xterm host	In-browser terminal surface (<code>TerminalHost.vue</code>) attached to a backend PTY session via <code>WS /api/workspaces/{workspace_id}/terminal</code> . Runs real shell commands in a workspace-scoped directory under <code>.local/workspaces/</code> (override with <code>AXON_WATCH_WORKSPACE_ROOT</code>).
Workspace	Logical operator context keyed by <code>workspace_id</code> . Today the API returns IDs only; rich catalog metadata is deferred.

Two repos, two apps

Do not confuse these:

	AXON-LOCAL (CURRENT AXON)	AXON-WATCH (AXON-X)
Default URL	<code>http://127.0.0.1:7734</code> (fallback)	<code>http://127.0.0.1:4173</code> (production operator)
Start command	<code>./start.sh</code> from axon-local	<code>./scripts/dev/up.sh</code> from axon-watch
Status	Legacy daily-driver / fallback	Primary operator console (v1)
Relationship	Source of parity targets and frozen plans	Implementation target for modernization

08 What Axon-X Is

Axon-X is the next-generation Axon console and operator environment.

The code lives in the `axon-watch` repo folder for now. That folder name is legacy/internal; the product name is **Axon-X**.

The target product combines:

- an IDE-style shell
- a control plane
- a dedicated watcher service

The implementation repo is here:

- `/run/media/vaxon/axon-data/repos/axon-nvme/repos/axon-watch`

The frozen planning source-of-truth still lives here:

- `/run/media/vaxon/axon-data/repos/axon-nvme/repos/axon-local/Plans/Axon-Watch/`

Important rule:

- read plans from `axon-local`
- implement in `axon-watch`

Do not casually move back and forth inventing new semantics in both places.

09 What State The Repo Is In Right Now

The repo is in an early but real bootstrap state.

What already exists:

- repo scaffold
- `console-web` shell skeleton
- `control-plane` FastAPI bootstrap service
- `axon-watch` FastAPI bootstrap service
- shared contract package
- verification harness
- health and readiness scripts

What is already real:

- root workspace scripts
- service startup/shutdown flow
- `/api/runtime/summary`, `/api/inbox`, `/api/runs`, and `/api/briefing` routes in the control plane
- shared contract types under `packages/shared-types/`
- shell consumption of canonical runtime summary, inbox, and run DTOs
- thin Monaco and xterm host surfaces in `console-web`

What is still intentionally thin or deferred:

- full signal production and deeper ranking beyond current inbox rule stack
- deep watch summary logic
- performance evidence for all budgets

What is now real in the thin slice (verified 2026-07-04):

- run create/complete/stop/resume lifecycle
- explicit approval boundary (`requires_approval`, approve/reject)
- review-ready entry, completion, and follow-up resume path
- SQLite-backed run persistence (survives control-plane restart)
- operator briefing loaded at bootstrap and rendered in the right dock (`BriefingPanel.vue`)
- two watch-produced inbox signals with multi-factor ranking (severity, recency, unresolved duration via `created_at`, status, action-type, workspace priority)
- workspace list API and shell workspace selector (IDs only)

- the shell is split into `TopBar`, `LeftSidebar`, `CenterWorkbench`, `RightDock`, and `StatusBar` regions with mockup-shell chrome
- Monaco host bound to canonical DTO documents and **workspace files on disk** with a nested explorer tree, lazy file loading, new-file creation, and active-file rename
- backend PTY terminal attachment for the selected workspace (real shell I/O via WebSocket)
- workspace-scoped conversation rehydration: `GET /api/workspaces/{workspace_id}/chat/thread` plus existing thread history read reloads the Conversation seam after page refresh. When no thread exists yet, the lookup returns HTTP 200 with null `thread_id` (not 404).

Workspace IDs (operator vs catalog):

- **Operator shell** uses `MOCKUP_WORKSPACE_IDS` only (`workspace_smoke`, `workspace_recsys`, ...). `mergeMockupWorkspaceCatalog()` trims API extras so `currentWorkspace` is always sidebar-visible.
- **Control-plane catalog** may still list fixture defaults (`workspace_alpha`, `workspace_bootstrap`) and run/inbox IDs for tests and API consumers; the shell does not select those as `currentWorkspace`.
- Bootstrap picks workspace deterministically: active run workspace (when visible) → `workspace_smoke` default → first mockup workspace.

Manual acceptance for reload-safe chat (use `workspace_smoke`):

1. `./scripts/dev/up.sh`
2. Open `http://127.0.0.1:4173`
3. Confirm `workspace_smoke` is selected (or select it)
4. Post a command in the Command seam
5. Hard reload the page
6. Conversation should rehydrate automatically for the same workspace when an active run or default bootstrap applies; if needed, re-select `workspace_smoke`

API-only check (any valid catalog ID):

```
curl -s http://127.0.0.1:8787/api/workspaces/workspace_smoke/chat/thread
curl -s http://127.0.0.1:8787/api/chat/threads/<thread_id>/history
```

What is **not** real yet despite similar-sounding names:

- **Full KAIRO operator presence** — the current shell has visual scaffolding, but spoken alerts, persona settings, and richer operator-presence behavior are still planned in axon-local docs
- **Full parity with axon-local** — intentional; see parity ledger for gaps

So this repo is not a fake mockup, but it is also not feature-complete or a drop-in replacement for the current Axon UI.

10 Source Of Truth Rules

When you are unsure what something should mean, check the planning bundle.

The most important frozen planning docs are:

- `PRODUCT.md`
- `ARCHITECTURE.md`
- `UI_SPEC.md`
- `UI_COMPOSITION_SPEC.md`
- `UI_VISUAL_DIRECTION.md`
- `UI_REFERENCE_ARCHETYPES.md`
- `run-state.md`
- `runtime-summary.md`
- `signal-events.md`
- `watch-api.md`
- `control-api.md`
- `PARITY_LEDGER.md`
- `FITNESS_FUNCTIONS.md`
- `TRANSITION_ARCHITECTURE.md`
- `CONTRACT_TESTING_SPEC.md`

Things you must not redefine casually

These are serial-only semantics and should not drift:

- canonical run phases
- runtime summary vocabulary
- signal identity, severity, and status
- approval semantics
- transition seams and rollback rules
- parity acceptance rules

If the plan seems wrong or insufficient:

- do not silently change it
- raise a proposed amendment instead

11 Repo Layout

Top-level structure:

```
axon-watch/  
  apps/  
    console-web/  
  services/  
    control-plane/  
    axon-watch/  
  packages/  
    shared-types/  
  docs/  
  scripts/  
    dev/  
    verify/  
  infra/  
  tests/
```

What each part owns

`apps/console-web/`

Owns the integrated shell:

- topbar
- left sidebar
- center workbench (including the embedded terminal dock)
- right dock
- status bar

It should consume canonical contracts, not invent backend truth.

`services/control-plane/`

Owns interactive backend responsibilities such as:

- UI-facing health/readiness
- runtime summary assembly
- later run-state, approvals, and workspace orchestration

`services/axon-watch/`

Owns watcher responsibilities such as:

- monitoring loops
- watch health/readiness

- later signal production
- later inbox/signal summaries

`packages/shared-types/`

Owns canonical shared contract types.

This is one of the most important directories in the repo.

If frontend and backend disagree about what a thing means, the fix should start here, not in ad hoc local types.

`scripts/dev/`

Owns local run helpers:

- start
- stop
- health checks

`scripts/verify/`

Owns cheap, repeatable verification:

- dependency direction checks
- DTO size checks
- ADR governance checks
- latency-budget check scaffolding

12 The Current Architecture In Plain English

Think of the new app as three cooperating parts:

1. the shell the user sees
2. the control plane that serves user-facing APIs
3. the watcher service that will eventually observe and normalize signals

The current implementation is following a thin-slice path:

- bootstrap the shell and services first
- land canonical contracts early
- add real endpoints against those contracts
- then deepen behavior

This is deliberate.

It avoids:

- giant rewrites
- hidden semantic drift
- UI and backend inventing different meanings

13 Detailed setup (first install)

1. Go to the repo

```
cd /run/media/vaxon/axon-data/repos/axon-nvme/repos/axon-watch
```

2. Install JavaScript dependencies

```
npm install
```

This is important because root workspace commands now rely on npm workspaces.

3. Review environment defaults

Default values are provided in:

- `.env.example`

If you need custom ports or local paths:

```
cp .env.example .env
```

Then edit `.env` as needed.

4. Start the local stack

```
./scripts/dev/up.sh
```

`up.sh` does not report success until all three services are actually reachable.

This starts:

- `console-web`
- `control-plane`
- `axon-watch`

Important startup contract:

- ports are fixed by `.env` / `.env.example`
- startup fails fast if `4173`, `8787`, or `8788` is already in use
- startup rolls back partial processes if readiness fails
- services stay detached after the launching shell exits

5. Check health

```
./scripts/dev/check-health.sh
```

Expected endpoints:

- console web: `http://127.0.0.1:4173`
- control plane health: `http://127.0.0.1:8787/api/health`
- watch health: `http://127.0.0.1:8788/internal/watch/health`
- briefing: `http://127.0.0.1:8787/api/briefing` (also loaded by the shell at bootstrap)
- runtime summary: `http://127.0.0.1:8787/api/runtime/summary`
- inbox: `http://127.0.0.1:8787/api/inbox`
- runs: `http://127.0.0.1:8787/api/runs`
- workspaces: `http://127.0.0.1:8787/api/workspaces`

`check-health.sh` also probes `/api/workspaces`.

6. Stop the stack

```
./scripts/dev/down.sh
```

14 What The Current App Does On Boot

Right now, the shell boots and loads five control-plane seams in parallel.

That flow looks like this:

1. `apps/console-web/src/main.ts` creates the Vue app
2. the shell store initializes
3. `loadBootstrapData()` is called
4. the frontend fetches `/api/runtime/summary`, `/api/inbox`, and `/api/briefing`
5. workspaces and runs load sequentially; `resolveBootstrapWorkspaceId()` sets `currentWorkspace` (active run workspace when sidebar-visible, else `workspace_smoke`)
6. workspace files and chat thread history load for that workspace
7. the shell renders topbar context, workspace state, editor/terminal workbench, right-dock seams, and status bar truth strips

Important limitations:

- **Operator workspace list** — shell and sidebar share the same `MOCKUP_WORKSPACE_IDS` catalog; catalog-only IDs such as `workspace_alpha` remain API-visible but are not selected in the shell
- **Bootstrap workspace selection** — deterministic via `resolveBootstrapWorkspaceId()` after sequential workspace + run load (no parallel race)
- **Chat rehydration** — scoped per workspace; messages posted under one ID do not appear when another workspace is selected; only the latest thread per workspace is returned
- **Chat orchestration** — `POST /api/chat/messages` returns operator + system + agent messages; new dispatches run bounded executor (`health` / `list files` / `read ...`) then `review_ready` with `command_execution` receipt
- **Workspace catalog** — sidebar uses seven mockup IDs; API may expose more — see `docs/WORKSPACE_CATALOG.md`
- Operator mode renders the right dock as **Run** → **Approvals** → **Signals** → **Conversation** → **Command** → **KAIRO Briefing** with the briefing card anchored at the bottom of the dock
- briefing data is projected across the dock rather than shown as one raw DTO dump: approvals stay in the approvals seam, top signals stay in the signals seam, and the KAIRO card stays summary/CTA-oriented
- Monaco host loads workspace files on disk and read-only DTO overview tabs
- xterm host attaches to a backend PTY scoped to the selected workspace directory
- inbox ranking uses severity, recency, unresolved duration (`created_at`), status, action-type, and a thin watch-owned workspace priority map

That is okay for this stage.

15 Locked Shell Layout

Locked 2026-07-04 — see `docs/UI_LAYOUT_LOCK.md` and `docs/adr/ADR-004-locked-console-shell-layout.md`.

Do not rearrange shell regions or dock seam order without a superseding ADR. Frozen planning in `axon-local/Plans/Axon-Watch/UI_COMPOSITION_SPEC.md` is amended to match this geometry.

Five-region grid:

REGION	COMPONENT	NOTES
Top bar	<code>TopBar.vue</code>	identity zone + runtime strip + KAIRO module + mode toggle
Left sidebar	<code>LeftSidebar.vue</code>	workspaces, optional explorer (IDE), status card
Center workbench	<code>CenterWorkbench.vue</code>	Monaco editor + embedded resizable terminal dock
Right dock	<code>RightDock.vue</code>	Run → Approvals → Signals (upper stack), KAIRO Briefing bottom hero
Status bar	<code>StatusBar.vue</code>	HUD runtime strip

KAIRO Briefing height tracks the workbench terminal dock via `--briefing-dock-height`.

16 Current Shell Layout

The locked shell matches the mockup/live screenshots:

- `TopBar` — brand frame, mockup breadcrumb/version context panel, DTO runtime strip, KAIRO presence module, Operator/IDE toggle, settings action
- `LeftSidebar` — workspace list first, workspace status card second, explorer tree only in IDE mode
- `CenterWorkbench` — editor tabbar + breadcrumb + Monaco editor above a resizable bottom terminal/log dock
- `RightDock` — run seam, approvals seam, signals seam, KAIRO briefing card
- `StatusBar` — persistent watch / run phase / signals / operator strip

This is the layout you should compare against the mockup and live screenshots, not the older single-file shell description from earlier thin slices.

17 Mockup / Live Parity Notes

Parity observations from the current mockup and live screenshots:

- the region geometry now largely matches the mockup: topbar, workspace rail, editor-over-terminal workbench, right dock, and bottom status strip
- the **runtime strip**, status bar zones, and workspace status card derive from live shell state / `RuntimeSummary`, but the topbar breadcrumb and runtime version chips are still mockup-style presentation helpers
- the sidebar and shell store share the same operator workspace catalog (`MOCKUP_WORKSPACE_IDS`)
- the dock uses operator-facing seam titles (`Active Run` , `Approvals` , `Signals` , `Conversation`) from `dock-seam-layout.ts`
- live shell refresh uses `GET /api/live/events` (SSE refresh hints) via `live-events-session.ts`, with visibility-aware polling fallback when EventSource is unavailable
- Operator mode renders the right dock as **Run** → **Approvals** → **Signals** → **Conversation** → **Command** → **KAIRO Briefing** with the briefing card anchored at the bottom of the dock
- the workbench terminal dock default height is ~240px (responsive cap 280px) unless the operator resizes it; height persists in session storage when customized

18 The Most Important Files Right Now

If you need to understand the current implementation quickly, read these first:

Shared contracts

- `packages/shared-types/src/index.ts`
- `packages/shared-types/src/run.ts`
- `packages/shared-types/src/runtime.ts`
- `packages/shared-types/src/signals.ts`
- `packages/shared-types/src/control-plane.ts`
- `packages/shared-types/src/watch.ts`

Fixture payloads

- `packages/shared-types/fixtures/runtime-summary.example.json`
- `packages/shared-types/fixtures/watch-summary.example.json`
- `packages/shared-types/fixtures/run-record.example.json`
- `packages/shared-types/fixtures/signal-event.example.json`

Control plane

- `services/control-plane/app/main.py`
- `services/control-plane/app/workspace_catalog.py`
- `services/control-plane/app/runs/service.py`
- `services/control-plane/app/runtime_summary_assembler.py`
- `services/control-plane/app/operator_briefing.py`

Frontend shell

- `apps/console-web/src/main.ts`
- `apps/console-web/src/api/control-plane.ts`
- `apps/console-web/src/stores/shell.ts`
- `apps/console-web/src/App.vue`
- `apps/console-web/src/components/EditorHost.vue`
- `apps/console-web/src/components/TerminalHost.vue`
- `apps/console-web/src/lib/workspace-documents.ts`

Verification

- `scripts/verify/README.md`
- `scripts/verify/verification_config.json`
- `tests/test_shared_contract_fixtures.py`
- `tests/test_control_plane_runtime_summary.py`
- `tests/test_verify_harness.py`

19 Verification Commands

Use these from the repo root.

CI, merge to `dev`, and employee agents: see [CI, merge, and worker agents](#).

Shared contract verification

```
npm run verify:shared-types
```

This confirms the shared-types package typechecks.

Contract verification

```
npm run verify:contracts
```

This verifies:

- shared contract package typing
- shared fixture tests
- control-plane run/inbox/runtime-summary/briefing behavior
- watch signal contract alignment

Full current verification bundle

```
npm run verify
```

This runs:

- contract verification
- console-web typecheck, unit test, and production build
- verify harness checks
- DTO size checks using representative fixtures

TEST-0 acceptance (`workspace_smoke`)

Requires the dev stack (`./scripts/dev/up.sh`).

```
npm run verify:test0
```

This runs, in order:

1. `./scripts/dev/check-health.sh`
2. Mission control unit tests (`operator-status-radar-view` , `workbench-terminal-split`)
3. Live acceptance — `tests/test_test0_workspace_smoke_acceptance.py` against `workspace_smoke`
(briefing Notice/Advise, git status, resume from review, inbox/runs)

4. `npm run verify`

See also `docs/OPERATOR_MISSION_CONTROL_V1.md` for the manual UI checklist.

21 Latency evidence (D1)

When the dev stack is running, collect warm-route timing samples and re-run verify with evidence files:

```
./scripts/dev/up.sh
npm run verify:evidence
npm run verify:nightly
```

This writes JSON under `.local/verify/` and passes them to `scripts/verify/all.py`.

Shell boot measurement:

- `scripts/dev/measure_shell_boot.py` records `shell_ready_ms`
- `auto` mode uses Playwright Chromium when installed; otherwise it measures the bootstrap critical path (index fetch + parallel `/api/runtime/summary`, `/api/inbox`, `/api/briefing`, `/api/workspaces`, `/api/runs`)
- optional full browser mode: `pip install playwright && playwright install chromium` then `AXON_WATCH_SHELL_BOOT_MODE=browser npm run verify:evidence`
- repo-local setup:

```
python3 -m venv .venv
. .venv/bin/activate
pip install -r requirements-dev.txt
python -m playwright install chromium
```

`collect-verify-evidence.sh` prefers `.venv/bin/python3` when present so `auto` mode can use Playwright.

Example fixture: `scripts/verify/fixtures/shell-boot-report.dev.json`.

`./scripts/dev/check-health.sh` also probes `GET /api/live/events` (SSE).

Frontend checks

```
npm run typecheck -w @axon-watch/console-web
npm run test -w @axon-watch/console-web
npm run build -w @axon-watch/console-web
```

Python syntax checks

```
python3 -m py_compile services/control-plane/app/main.py services/control-plane/app/runs/service.py services/control-plane/app/domain/run_state.py services/control-plane/app/domain/run_transitions.py services/control-plane/app/operator_briefing.py
```

Service tests

```
python3 -m unittest discover -s tests
```

22 What PASS / PENDING / FAIL Mean In Verification

The verify harness uses explicit states.

PASS

The supplied evidence met the rule.

Examples:

- DTO payload fits its size budget
- dependency direction rule is respected

PENDING

The check exists, but the slice does not provide enough real implementation or evidence yet.

Examples:

- shell boot timing evidence not supplied yet
- runtime summary latency budget not yet measured
- no numbered ADR yet exists

PENDING is not automatically bad at this stage.

It means the governance has been created ahead of the full implementation.

FAIL

The rule was violated or the harness itself is broken.

23 Common Working Patterns

When adding frontend work

Use real shared types from `packages/shared-types/`.

Do not:

- invent local copies of DTOs
- widen semantics casually in component code
- hide backend meaning in placeholder strings

Good pattern:

- define or refine contract in shared-types
- use it in store/API layer
- render it in Vue components

When adding control-plane work

Prefer:

- canonical DTO assembly
- light endpoint handlers
- explicit payload validation

Avoid:

- stuffing frontend-specific assumptions into the backend
- returning huge boot payloads
- redefining contract shapes inline

When adding watch work

Prefer:

- narrow, owned summaries
- canonical signal/event identities

- explicit watch responsibilities

Avoid:

- pulling UI logic into the watch service
- importing control-plane domain semantics directly
- inventing competing signal vocabulary

24 Troubleshooting

25 Debugging playbook

For failed employee agents, first use the cause-specific recovery table in [Stuck agent recovery & School Operations Phase 1](#). In particular, do not restart a healthy control plane just to clear a failed card: a restart interrupts in-flight shifts.

Use this order when something breaks:

1. **Stack health** — `axonhealth` (or `./scripts/dev/check-health.sh`).
2. **Empty shell / Runtime unavailable** — `axonrevive`, then hard-refresh `:4173`. Do **not** rely on `./scripts/dev/down.sh` when systemd owns the ports.
3. **Soft refresh** — `axonrestart` after backend route changes (if the API still answers).
4. **Required connector / tunnel / vault 503** — `axonfixconnectors` (see below). Local `axonhealth` can still be green while the **public** console probe is red.
5. **Stale UI on `:4173`** — rebuild console-web (`npm run build -w @axon-watch/console-web`), `systemctl --user restart console-web.service`, then hard-refresh. Source-only / `:5173` Vite edits do **not** update the systemd bundle.
6. **Browser cache** — hard refresh `:4173` (`Ctrl+Shift+R`) after console-web bundle changes.
7. **Connectors truth** — Mission Control → **Connectors** rail or `GET /api/connectors`.
8. **Gate scripts** — `npm run verify:production-operator`, then the slice gate (`verify:testN / verify:shell-commands`).
9. **Logs** — `journalctl --user -u control-plane.service -n 80` (always-on) or `.local/logs/` (dev bootstrap).

Symptom-specific fixes continue in the sections below.

Problem: REQUIRED CONNECTOR DOWN / tunnel token missing / Vault unlock HTTP 503

What you see

- Mission Control chip: **1 REQUIRED CONNECTOR DOWN, degraded, tunnel token missing**
- Connectors rail: **Console web** `REQUIRED` → unavailable / probe failed
- Cloudflare tunnel row: `tunnel token missing (auth=missing)`
- Vault page (`/vault`): unlock fails with red banner
`Watch vault API HTTP 503: AXON_WATCH_INTERNAL_SERVICE_TOKEN is required when the operator surface is remotely reachable`
- Status line may still say **WATCH CONNECTED** and `axonhealth` can still pass — because **local** `:4173` / `:8787` are fine.

Current rule: required `console_web` probes **loopback** (`AXON_WATCH_CONSOLE_WEB_BASE_URL` , usually `:4173`). Cloudflare/public reachability is optional `public_ingress` . A tunnel flap alone must not mark Mission Control degraded.

Older builds probed required `console_web` against `AXON_WATCH_PUBLIC_BASE_URL` (e.g. `https://axon.edudashpro.org.za/api/health`), which incorrectly tied local ONLINE to Cloudflare.

Why it happens (failure chain)

1. `AXON_WATCH_PUBLIC_BASE_URL` is a non-loopback hostname → deployment is **remotely reachable**.
2. Remotely reachable hosts **require** a shared `AXON_WATCH_INTERNAL_SERVICE_TOKEN` on both control-plane and axon-watch (same value in `~/.config/axon-watch/deployment.env`).
3. If that token is missing, control-plane → watch **mutating** calls (vault unlock, tunnel start) return **HTTP 503**.
4. Vault stays **locked** → tunnel token cannot resolve from encrypted vault secrets → managed `cloudflared` will not start → public health returns Cloudflare **1033** → optional `public_ingress` goes soft (local required connectors should still be green).
5. Auto-unlock keyfile may show as enabled on `/vault` , but **auto-unlock is refused by default** when remotely reachable. On this trusted always-on host set `AXON_WATCH_ALLOW_VAULT_AUTO_UNLOCK=1` in `~/.config/axon-watch/deployment.env` , then `axonrestart` (or Enable auto-unlock in `/vault`).
6. Separately: a slow `/api/vault/status` (Cursor/Codex CLI probes) used to exceed the console's **12s** fetch budget and show a false **SETUP REQUIRED** — probes are now short-timeout + cached.

Also see [docs/how-to/autonomy-gates-and-service-identity.md](#) and [docs/NATIVE_TUNNEL_CONTROL.md](#) .

Fix (copy-paste)

One-word path (preferred):


```
# 1) Only if vault unlock shows INTERNAL_SERVICE_TOKEN HTTP 503:
axonfixconnectors --ensure-internal-token --restart

# 2) Browser: /vault → master password + 2FA → Remember me → UNLOCK
#   Confirm secret name: cloudflare_tunnel_token
#   (aliases: AXON_CLOUDFLARE_TUNNEL_TOKEN / CLOUDFLARE_TUNNEL_TOKEN / TUNNEL_TOKEN)

# 3) Start tunnel + reprobe (reads AXON_WATCH_OPERATOR_TOKEN from deployment.env)
axonfixconnectors
```

Do not restart `axon-watch` after a successful vault unlock unless auto-unlock is permitted on this host (`AXON_WATCH_ALLOW_VAULT_AUTO_UNLOCK=1` + keyfile) or you are ready to unlock again. Vault unlock is **in-process** on the watch service; without auto-unlock a watch restart drops the session. Managed `cloudflared` can also stop with the watch unit.

Other facts for this host:

- Mutating CP routes (`POST /api/tunnel/start`, `reprobe_connector`) need `Authorization: Bearer <AXON_WATCH_OPERATOR_TOKEN>` under `local_token` (forced when the public URL is non-loopback). `axonfixconnectors` loads that token from `~/.config/axon-watch/deployment.env`. Mission Control can mutate via a desktop session cookie instead.
- Auto-unlock is refused by default while remotely reachable. Trusted always-on host override: `AXON_WATCH_ALLOW_VAULT_AUTO_UNLOCK=1` in `deployment.env`, then `axonrestart`. Confirm `/vault` no longer shows the remote-disable banner after Enable.
- Soft cutover (`remote=http://localhost:7734` via public-origin-proxy) is healthy but not final. To point Cloudflare directly at Axon-X `:4173`:

```
# Needs Cloudflare API token with Account → Cloudflare Tunnel → Edit
# Store as CF_API_TOKEN in deployment.env or vault, then:
./scripts/ops/set-tunnel-ingress-4173.sh
# Reprobe Cloudflare tunnel – expect ingress_matches_axon (no soft-cutover chip)
```

Voice says “open Runtime or vault” but the job failed for everyone

That line was a **misclassified Cursor usage-limit failure**. Lane B used to append “Open Runtime or `/vault`” on every CLI failure, so VAXON treated usage blocks as vault problems. Copy now distinguishes:

REAL CAUSE	SPOKEN / ROSTER NEXT STEP
<code>ActionRequiredError</code> / hit usage limit	Raise Cursor limit or switch model
Auth probe timeout	Check <code>cursor agent status</code>
Not signed in / vault locked	Login or unlock <code>/vault</code>

Guard-rails: continuous scheduler skips the whole workspace after any role hits usage limits (Cursor quota is account-wide). Pause Fleet Controls when burning quota on a known limit. - `cloudflare_tunnel` / `public_ingress` are **optional** (`required: false`). **REQUIRED CONNECTOR DOWN** is driven by required **loopback** probes — `control_plane` + local `console_web` . - Soft cutover is normal: remote ingress may still target `http://localhost:7734` while `axon-public-origin-proxy` forwards to `:4173` . Healthy tunnel detail: **active soft cutover** . - Connector probes must use a non-default User-Agent. Cloudflare returns **403** to stock `Python-urllib` , which used to surface as generic `probe failed` even when `curl` and the tunnel probe succeeded (fixed in connector probe headers).

Install / refresh the PATH wrapper if needed:

```
./scripts/ops/install-bin-wrappers.sh
# or without PATH:
./scripts/ops/axonfixconnectors.sh --ensure-internal-token --restart
```

Manual path (same outcome):

```
# Add shared internal token only if missing (do not commit this file)
echo "AXON_WATCH_INTERNAL_SERVICE_TOKEN=$(openssl rand -hex 24)" >> ~/.config/axon-
watch/deployment.env
axonrestart

# Unlock Vault in the UI (Remember me), then:
source ~/.config/axon-watch/deployment.env # or export OPERATOR token another way
curl -sS -X POST http://127.0.0.1:8787/api/tunnel/start \
  -H "Authorization: Bearer ${AXON_WATCH_OPERATOR_TOKEN}" | python3 -m json.tool
curl -sS -X POST http://127.0.0.1:8787/api/watch/commands \
  -H "Authorization: Bearer ${AXON_WATCH_OPERATOR_TOKEN}" \
  -H 'Content-Type: application/json' \
  -d
'{"command_type":"reprobe_connector","target_type":"connector","target_id":"console_web","re
quested_by":"operator"}'
```

Durable tunnel auth (survives watch restart): put the named-tunnel token in env (or keep unlocking after every watch restart):

```
# ~/.config/axon-watch/deployment.env
AXON_CLOUDFLARE_TUNNEL_TOKEN=<named-tunnel-token-from-cloudflare>
axonrestart
axonfixconnectors
```

Accepted vault/env names: `AXON_CLOUDFLARE_TUNNEL_TOKEN` , `CLOUDFLARE_TUNNEL_TOKEN` , `cloudflare_tunnel_token` , `TUNNEL_TOKEN` . When the vault secret `cloudflare_tunnel_token` is loaded, status may show `auth_source=settings` (resolver labels a non-empty stored value that way) — that still means a token is available.

Verify

```
axonfixconnectors          # required_unavailable should be 0
curl -sS http://127.0.0.1:8787/api/tunnel/status | python3 -m json.tool
curl -sS --max-time 15 https://axon.edudashpro.org.za/api/health
# hard-refresh Mission Control – degraded / REQUIRED CONNECTOR DOWN chips should clear
```

Problem: Runtime unavailable / No workspace selected / Briefing unavailable

This is the **wedged control-plane** pattern. The Vue shell is up, but `/api/*` times out (even `/api/health`). Bootstrap never selects a workspace, so explorer + chat look empty.

```
axonrevive
# hard-refresh http://127.0.0.1:4173
```

Why `./scripts/dev/down.sh` / `up.sh` fail here: this host runs **user systemd units** (`control-plane.service`, etc.). Dev down/up skip those listeners. Soft `systemctl --user restart` can also hang on a stuck worker — `axonrevive` force-kills first.

Problem: Vite reports `ECONNREFUSED 127.0.0.1:8787`

Use `scripts/ops/run-5173.sh`; follow the [restart playbook](#).

Problem: online research falls back / Google search returns 403

See [SearXNG research](#) for provider order and Google 403 recovery.

Problem: `./scripts/dev/up.sh` fails or the frontend does not start

Check:

1. Did you run `npm install` at repo root?
2. Are ports `4173`, `8787`, and `8788` free?
3. Did `up.sh` print a specific port conflict or readiness failure message?

Current expected path:

- root install: `npm install`
- root startup: `./scripts/dev/up.sh`
- health check: `./scripts/dev/check-health.sh`

If startup fails, inspect:

- `.local/logs/console-web.log`

- `.local/logs/control-plane.log`
- `.local/logs/axon-watch.log`

Problem: `check-health.sh` fails

This usually means one or more services never started.

Steps:

1. stop everything:

```
./scripts/dev/down.sh
```

2. start again:

```
./scripts/dev/up.sh
```

3. inspect logs:

- `.local/logs/control-plane.log`
- `.local/logs/axon-watch.log`
- `.local/logs/console-web.log`

4. re-run:

```
./scripts/dev/check-health.sh
```

Problem: stale pid files block startup

`up.sh` now clears stale pid files automatically before it checks for a live stack.

First try:

```
./scripts/dev/down.sh
```

If that is not enough, inspect:

- `.local/pids/`
- `.local/logs/console-web.log`
- `.local/logs/control-plane.log`
- `.local/logs/axon-watch.log`

If `up.sh` still fails, it usually means one of the configured ports is held by another live process. Free the port or stop the other process before retrying.

Problem: the shell loads but runtime summary is unavailable

Check:

1. Is control-plane running?
2. Does this endpoint work?

```
curl -fsS http://127.0.0.1:8787/api/runtime/summary | python3 -m json.tool
```

3. If running through Vite dev server, is `/api` proxied to control-plane?

Current dev proxy lives in:

- `apps/console-web/vite.config.ts`

You can also override the target with:

- `VITE_CONTROL_PLANE_BASE_URL`

Problem: `npm run verify` shows `PENDING`

That is often expected at this stage.

Examples of currently expected pending checks:

- shell boot readiness
- runtime summary latency
- watch summary latency
- numbered ADR presence

Treat `PENDING` as:

- “scaffold exists, evidence not landed yet”

Treat `FAIL` as:

- “something is wrong right now”

Problem: dependency direction checks fail

Check for forbidden imports across boundaries.

The current verify harness guards things like:

- watch service must not depend on UI internals
- frontend must not depend on watch internals
- future domain-layer cross-imports must stay clean

If a dependency check fails:

1. remove the cross-boundary import
2. move shared meaning into `packages/shared-types/` if appropriate
3. use an adapter or API boundary instead of direct reach-through

Problem: contract tests fail after changing fixtures or DTOs

This usually means one of three things:

1. you changed a canonical field name
2. you widened/narrowed a type without updating the fixture
3. you changed semantics without updating the contract layer first

Fix sequence:

1. check the frozen planning doc
2. check the shared contract type
3. check the fixture
4. check the consumer/provider test

Do not patch the frontend alone and ignore the canonical contract.

Problem: you are unsure whether something should become a new ADR

Ask:

1. Is this a real architecture or process decision?
2. Is it broader than one small implementation detail?
3. Would future readers need to know why this choice was made?

If yes, use ADR governance.

Current ADR process docs live in:

- `docs/adr/README.md`
- `docs/adr/_TEMPLATE.md`

Note:

- accepted ADRs should be numbered

- accepted ADRs should not be rewritten materially
- changed decisions should be superseded, not silently edited

26 Tips And Tricks

27 Snippet cookbook

One-word stack commands (always-on host)

Installed on PATH via `~/local/bin` → repo `bin/`:

COMMAND	EXPANDS TO	WHEN TO USE
<code>axonhealth</code>	<code>./scripts/dev/check-health.sh</code>	Quick “is the stack OK?”
<code>axonrestart</code>	<code>systemctl --user restart axon-watch + control-plane + console-web</code> , then health	Soft restart when APIs still respond
<code>axonrevive</code>	Force-kill control-plane → restart all three → health	Empty shell, hung health, wedged worker
<code>axonfixconnectors</code>	Diagnose tunnel/connectors; optional <code>--ensure-internal-token --restart</code>	REQUIRED CONNECTOR DOWN, tunnel token missing, vault unlock HTTP 503

```
axonhealth
axonrestart
axonrevive
axonfixconnectors --ensure-internal-token --restart # vault 503 / missing internal token
axonfixconnectors # after vault unlock: start tunnel +
reprobe
```

Repo scripts (same behavior without PATH):

```
./scripts/ops/axonhealth.sh
./scripts/ops/axonrestart.sh
./scripts/ops/axonrevive.sh
./scripts/ops/axonfixconnectors.sh
```

Open console after revive: **http://127.0.0.1:4173** (hard-refresh).

Full connector/tunnel recovery write-up: [Problem: REQUIRED CONNECTOR DOWN](#).

Console-web rebuild (always-on :4173)

:4173 serves the **built** `apps/console-web/dist` via systemd `console-web.service`. Source edits (including Vite :5173) are **not** live on :4173 until:

```
npm run build -w @axon-watch/console-web
systemctl --user restart console-web.service
# then hard-refresh http://127.0.0.1:4173
```

Local verify loop

```
./scripts/ops/change-verify-loop.sh          # dirty working tree
./scripts/ops/change-verify-loop.sh --head-only # committed HEAD only
./scripts/ops/change-verify-loop.sh --watch
```

Dev bootstrap (when systemd is not owning ports)

```
./scripts/dev/up.sh
./scripts/dev/down.sh
./scripts/dev/check-health.sh
```

28 Upgrading and updating

After pulling new commits or changing dependencies:

```
cd /run/media/vaxon/axon-data/repos/axon-nvme/repos/axon-watch
git pull
npm install
./scripts/dev/down.sh && ./scripts/dev/up.sh
npm run verify:production-operator
```

Sync planning docs back to axon-local when Axon-Watch planning changed:

```
python3 scripts/ops/sync_planning_mirror_to_axon_local.py
```

Hard-refresh :4173 after frontend changes. Restart the stack after control-plane or watch route changes.

Tip 1: Read the planning bundle before expanding semantics

The planning bundle is not optional context. It is the definition of intended meaning.

Tip 2: Treat `packages/shared-types/` as sacred

If the frontend and backend disagree, fix it here first.

Tip 3: Keep slices narrow

The repo is intentionally growing by thin slices.

If you are touching:

- `run-state`
- `runtime summary`
- `signals`
- `approvals`

then keep the slice tightly bounded and verifiable.

Tip 4: Prefer fixtures early

Fixtures are useful in early slices because they:

- keep contracts concrete
- make DTO size checks easy
- let frontend and backend agree before deeper logic exists

Tip 5: Use the verify harness even when it is not strict

`PENDING` today becomes `PASS` or `FAIL` tomorrow.

The harness is part of how this repo avoids drifting back into a monolith.

Tip 6: Distinguish “bootstrap-real” from “feature-real”

Something can be real enough to run and still be intentionally shallow.

Current examples:

- `runtime summary endpoint` is real
- `runtime summary assembly` is still bootstrap-thin
- `axon-watch` emits `signal_runtime_summary_degraded` with bootstrap-aware copy (`Bootstrap: runtime summary stale`) while watch connectivity is healthy — this is expected local scaffolding, not a production

outage

That distinction matters during review.

Tip 7: Do not overreact to incomplete polish

Prefer contracts and verify harness first; cosmetic cleanup can wait.

29 What A Good Next Slice Looks Like

A good next slice should:

- keep shared contract semantics stable
- improve one owned behavior
- preserve boot simplicity
- come with verification

Completed thin slices (do not re-do):

- bootstrap + shared contracts + runtime summary
- first watch signal path
- npm workspace dev ergonomics
- first run lifecycle (create → executing → complete)
- startup supervision reliability (`scripts/dev/lib/common.sh`)
- stop/resume, approval, review-ready, SQLite persistence, and briefing-backed dock projections
- workspace list + backend PTY terminal + file-backed Monaco host + nested explorer tree + new-file creation + active-file rename + resizable bottom terminal dock
- richer inbox ranking (severity, recency, unresolved duration, status, action-type, workspace priority)
- split shell regions (`TopBar` , `LeftSidebar` , `CenterWorkbench` , `RightDock` , `StatusBar`) with mockup-shell HUD chrome
- Conversation and Command dock seams backed by control-plane chat endpoints (`POST /api/chat/messages` , `GET /api/workspaces/{workspace_id}/chat/thread` , `GET /api/chat/threads/{thread_id}/history`)

Suggested next slices (2026-07-04):

1. **Coordinator** — KAIRO operator-presence integration when explicitly assigned
2. **Lane B** — agent orchestration hook for chat messages (beyond system ack stub)

Bad next slices:

- broad UI rewrite
- expanding multiple semantic families at once
- changing run-state and signal-state in one uncontrolled pass
- skipping verification because “it is still early”
- claiming full IDE parity when workspace file operations are still intentionally thin-slice (open, edit, save, create, rename) rather than a full VS Code clone

30 Final Guidance

If you are unsure what to do next, choose the smaller move that:

- preserves ownership
- strengthens verification
- reduces placeholders
- keeps the shell boot-safe

That is the design center of this repo right now.

31 CI, merge workflow, and worker agents

Updated: 2026-07-30

This chapter explains how Axon-X lands code through GitHub CI, how that relates to the `dev` branch, and how **company employee agents** (the roster in IDE / Mission Control) work — including how to tell whether they are actually running and how good the automation is today.

It also covers **DashPro self-hosted CI** (plain English + RAM) and why **canary OTA** is not an “EAS Build” in the Expo UI.

Related: [docs/CI_GATES.md](#) , [docs/planning/EXECUTION_PLAN.md](#) , [recent-operator-features.md](#) , [docs/planning/DASHPRO_CI_AGENT_PLAYBOOK.md](#)

Names — do not confuse these

NAME	WHAT IT IS
Axon-X Fast Gate	GitHub Actions workflow (<code>.github/workflows/fast-gate.yml</code>) — automated tests on every push and PR
GitHub → Agents tab	GitHub product UI (Copilot / cloud agents). Not the Axon company roster
Company roster / employee agents	Config-driven roles (<code>Night Watch</code> , <code>Shell Craft</code> , ...) in <code>config/workspace-agents.json</code> — continuous workers on a schedule
Lane B / IDE agent	Cursor CLI agent runs you start from the IDE composer (operator-driven)
Cursor worker-server	Internal Cursor subprocess; not an Axon employee by itself

CI and merge — how it works

Integration truth

BRANCH	ROLE
<code>dev</code>	Integration branch. Protected on GitHub — <code>fast-gate</code> must pass before merge
<code>feat/*</code>	One feature slice per branch (e.g. <code>feat/operator-console</code>)
<code>master</code>	Legacy default; not the day-to-day Axon-X integration target

Day-to-day console work targets `dev` via pull request. Do not camp on `master` for operator-console slices.

What triggers CI

Workflow: **Axon-X Fast Gate** (`.github/workflows/fast-gate.yml`)

Runs on:

- every **push** to any branch
- every **pull_request** (open, sync, reopen)

In the GitHub UI: **Actions** → workflow runs show green ✓ or red ✗. A PR also shows the same check on the PR page.

Poll Fast Gate from the terminal (required after every push)

Agents and operators should **not** wait for someone to notice a red run:

```
./scripts/ops/watch-fast-gate.sh
## or follow a known run id:
./scripts/ops/watch-fast-gate.sh 29925529734
```

On failure:

```
gh run view --log-failed
## Typical Fast Gate early fail: file-size ratchets in
scripts/guardrails/hotspot_budgets.json
```

Axon-X company ownership: Rowan (watcher) + Mira (Lead) own Fast Gate for `workspace_axon_watch`. Quinn (integrations) supports Actions wiring.

How to know older red runs are “fixed”

GitHub **never** flips a finished red run green. Older failures on the same branch stay red forever. The only merge signal that matters is the **tip of the branch**:

1. `git rev-parse HEAD` matches the green Fast Gate run's `headSha`.
2. `gh run list --branch <branch> --workflow "Axon-X Fast Gate" --limit 1` shows `conclusion=success`.
3. Or Actions UI: newest Fast Gate on that branch is green for the same commit.

Typical early fails: hard **file-size ratchets**, planning **manifest hash** drift, or missing **workspace bindings**. Fix tip → push → watch Fast Gate. Uncommitted local WIP is not CI-proven until it is committed, pushed, and green.

What Fast Gate runs (~2–3 minutes)

1. `npm run verify:contracts` — shared types, control-plane/watch contract unit tests, file-size ratchets
2. `npm run verify:hotspot-changes` — critical hotspots must shrink or stay within budget when changed (PR compares against base branch)
3. `npm run verify:console-web` — Vue typecheck, full Vitest suite, production build
4. `scripts/verify/all.py --strict-pending` — import boundaries, DTO budgets, ADR governance, latency scaffold

If any step fails, the run is red and `dev` **cannot merge** that PR until fixed.

Operator merge workflow

```
feat/my-slice → push → open PR to dev → fast-gate green → merge → git pull origin dev
```

Step by step:

1. **Work on a feature branch** (not directly on `dev`).

2. **Commit focused slices** — one concern per PR when possible (worker scheduler separate from IDE chrome, separate from connector parity).
3. **Push** the branch (`git push -u origin feat/my-slice`).
4. **Open a PR** to `dev` (`gh pr create --base dev --head feat/my-slice`).
5. **Wait for** `fast-gate` — green on the PR (not only on an earlier push).
6. **Merge** when green (GitHub UI or `gh pr merge`).
7. **Sync locally:** `git fetch origin dev && git checkout feat/my-slice && git merge origin/dev` (or rebase per team habit).

Important: A green check on an old push does **not** guarantee the open PR is green. Always read the check on the PR itself.

Local preflight (before you push)

```
npm run verify:preflight # via scripts/sc on commit – blocks bad commits locally
npm run verify:contracts # fastest backend + ratchet signal
cd apps/console-web && npm run typecheck && npm run test
git status --short      # must contain no source changes or local runtime artifacts
```

Before opening a PR, fetch `dev`, integrate it on the feature branch, rerun the preflight, and inspect the complete branch diff:

```
git fetch origin dev
git merge origin/dev
git diff --check
git diff --stat origin/dev...HEAD
```

Do not merge the feature branch directly into protected `dev`, force-push, or hide a failing check. Local runtime artifacts (`control-plane.sqlite3`, debug logs, shell completion dumps) are not source and must remain ignored.

Full local bundle (slower):

```
npm run verify
```

Nightly gate (not on every PR)

Workflow: `.github/workflows/nightly-verify.yml` — live stack evidence. See [docs/CI_GATES.md](#).

Worker / employee agents — how they work

What they are

Axon-X models each workspace as a **small company** with named employees (lead, watcher, frontend, backend, integrations). Configuration lives in:

- `config/workspace-agents.json` — companies, roles, schedules, `owns` text
- `services/control-plane/app/workspace_agents/` — load config, derive status, scheduler

UI:

- **IDE** → **TEAM tab** (`IdeTeamPanel` → `CompanyRosterPanel`)
- **Talk / Assign** — pre-fills the IDE composer (operator-driven); does not by itself start a scheduled worker

Schedules

SCHEDULE	MEANING
<code>on_demand</code>	No automatic shift (typical for lead)
<code>always_on</code>	Watcher-style; scheduler starts bounded shifts (e.g. Night Watch)
<code>continuous</code>	Specialist roles (frontend, backend, integrations) get recurring shifts

Skipped from auto-schedule: `lead` , `overview_agent` .

Scheduler loop (WA-1)

When the control plane starts (`services/control-plane/app/bootstrap.py`):

1. `start_continuous_worker_scheduler()` runs a background tick (default **every 45s**).
2. Each tick, for each `always_on` / `continuous` employee without an active role-tagged run, the scheduler: -
`create_run(..., employee_role=<role>)` — run record in SQLite -
`dispatch_continuous_worker_run()` — headless Lane B with a role-scoped prompt (`worker_prompt.py`)
→ `worker_dispatch.py` → Cursor/Codex CLI)

Guards (so one restart cannot flood the fleet):

- `MAX_STARTS_PER_TICK` (6) — cap new starts per tick
- `MAX_ACTIVE_EXECUTING` (24) — skip new starts when executing debt is high
- Per-role gate — only one busy shift per (`workspace_id` , `employee_role`) at a time

Environment knobs (`.env` / vault)

VARIABLE	DEFAULT	PURPOSE
<code>AXON_WATCH_WORKER_SCHEDULER</code>	<code>1</code>	Enable/disable scheduler
<code>AXON_WATCH_WORKER_SCHEDULER_DISPATCH</code>	<code>1</code>	Lane B dispatch after <code>create_run</code>
<code>AXON_WATCH_WORKER_SCHEDULER_INTERVAL_SECONDS</code>	<code>45</code>	Tick interval
<code>AXON_WATCH_WORKER_RUN_STALE_SECONDS</code>	<code>720</code>	Fail/cancel idle role-tagged shifts older than this
<code>AXON_WATCH_LEAD_RUN_STALE_SECONDS</code>	<code>1800</code>	Lead-role idle TTL (defaults to at least 30m so fan-out / board work survives)
<code>AXON_WATCH_EMPLOYEE_RUN_RETENTION_PER_ROLE</code>	<code>8</code>	Keep this many finished role-tagged runs per workspace/role
<code>AXON_WATCH_REVIEW_READY_STALE_SECONDS</code>	<code>14400</code>	Auto-complete idle untagged <code>review_ready</code> checkpoints older than this

Disable scheduler in tests: `AXON_WATCH_WORKER_SCHEDULER=0`.

Staffing file override: `AXON_WATCH_WORKSPACE_AGENTS_FILE` (path to JSON).

Stale shifts and history retention

Hung role-tagged runs can block the next shift for that role. The control plane:

- Reaps idle employee runs on each scheduler tick and on startup
- Exposes `POST /api/runs/reconcile-stale` for an on-demand reap
- Drains finished employee-run history beyond the per-role retention window on startup / `POST /api/runs/prune-employee-history` (scheduler ticks prune in smaller batches so shifts stay responsive)
- Completes abandoned untagged `review_ready` checkpoints on startup, each scheduler tick, and `POST /api/runs/reconcile-review-ready` so briefing stays honest when review was never dismissed

Approval waits are never auto-completed or pruned. Fresh review checkpoints stay until the idle TTL elapses.

How to know if employee agents are working

1. Company roster (UI)

IDE mode → **TEAM** tab (or operator company panel). Each employee shows a status: `idle`, `watching`, `executing`, `planning`, `verifying`, etc.

Good sign: specialists (`frontend` , `backend` , ...) move to `executing` when their shift is active, not only the lead mirroring any workspace run.

2. Mission Control / runs API

```
curl -sS http://127.0.0.1:8787/api/runs | python3 -c "  
import sys, json  
runs = json.load(sys.stdin)  
if isinstance(runs, dict): runs = runs.get('runs', runs.get('items', []))  
for r in runs:  
    if r.get('employee_role') and r.get('phase') not in ('completed','failed','cancelled'):  
        print(r['run_id'], r['workspace_id'], r['employee_role'], r['phase'],  
r.get('summary','')[:50])  
"
```

Good sign:

- `employee_role` is set (`watcher` , `frontend` , ...)
- `phase` is `executing` (or progresses to `completed` / `failed` with receipts)
- Summary mentions `continuous worker shift`

Bad sign (phantom worker): `executing` forever with **no** `cursor-agent` process and **no** `runtime_dispatch` receipt on the run history.

When Team says “last job failed” (verify before retrying)

The red Team badge is a **roster projection**, not the authoritative run result. Before tapping **Try again** on the agent review strip (above the composer), compare these three surfaces:

```
## 1. Current roster projection  
curl -sS http://127.0.0.1:8787/api/workspaces/workspace_axon_watch/company \  
| python3 -m json.tool  
  
## 2. Authoritative run record  
curl -sS http://127.0.0.1:8787/api/runs/<run_id> | python3 -m json.tool  
  
## 3. Transition/receipt history (the reason, not only the red label)  
curl -sS http://127.0.0.1:8787/api/runs/<run_id>/history \  
| python3 -m json.tool
```

Interpret the result:

- `phase=completed` plus `operator_complete` is success. Do **not** retry it because an older red badge or a sentence inside the transcript says “failed”.
- `phase=failed` plus `run_failed` is a real run failure. Read its receipt.
- `finalization_error` after `run_failed` is usually a **late-result race**: the scheduler failed the run first; Cursor returned later and could no longer complete a run already in `failed` .

- A Lead rollout saying `0 completed · 2 failed` can mean **two specialist assignments failed**, not that the Lead process failed twice.
- `missing or failing acceptance_evidence (Gate 6)` means the specialist did work but did not deliver a passing acceptance receipt.
- `Critical Review Clause missing` means the final answer omitted the required terminal `Confidence: N/10` line. This is a delivery-contract failure, not proof that the implementation itself was wrong.
- If roster says failed but the newest role run is completed, refresh/reopen Team and wait for the next company poll. Do not create another retry loop.

Mira incident receipt — 29 Jul 2026

What looked like “Mira failed twice” was three different facts:

1. Lead plan `lead-plan-78bfcac685b74588` had **two failed specialist tasks** (Quinn/integrations and Rowan/watcher). Both exhausted their three-attempt budgets. The surfaced blockers were Gate 6 acceptance evidence and the missing Confidence clause.
2. Mira retry runs `run_dc05055ef75f` and `run_fa7cc11f311e` did **not** fail: both reached `completed`; Critical Review recorded 7/10 and 8/10.
3. The real older Mira failure was `run_333574cdce66`: the scheduler marked it failed after `753s > 720s` without progress. Cursor returned about two minutes later, so finalization also logged `complete requires ... found failed`. That is one timed-out run with a late finalization error, not two independent implementation failures.

The safe operator response is: open run history, identify whether the failure belongs to Lead or a specialist, fix the named delivery contract, then retry **once**. Confirm the retry reaches `completed` before trying again.

3. Live CLI processes

```
pgrep -af 'cursor-agent.*agent --print' | head
```

You should see prompts like *"You are Shell Craft, the frontend employee for workspace workspace_axon_watch..."* when dispatch is active.

4. Control plane health

```
./scripts/dev/check-health.sh
curl -sS http://127.0.0.1:8787/api/health
```

5. Agent shift digests (optional)

When agents write them, receipts live under `docs/ops/agent-reports/` (untracked until committed). Use for human-readable "what this shift did."

6. Unit tests

```
PYTHONPATH=services/control-plane python3 -B -m unittest \
  tests.test_workspace_agent_scheduler \
  tests.test_workspace_worker_prompt \
  tests.test_workspace_agents -q
```

How good are they? (honest assessment)

DIMENSION	TODAY	TARGET
Visibility	Roster + runs show role and phase	Same, plus ranked coaching in briefing (SA-2)
Autonomy	Bounded scheduled shifts with role prompts	Continuous office cadence + checkpoint "meetings"
Continuity	One shift per role; tick every 45s	Stale reaper + handoff memory across shifts
Quality of work	Depends on Lane B + workspace scope; can edit/run in bound repo	Role-scoped playbooks (e.g. DashPro CI triage)
Safety	Debt caps, no auto-start for lead; dispatch uses controlled prompts	Autonomy ladder per workspace (<code>observe</code> → <code>full</code>)

They are real workers when: Lane B dispatch is on, `cursor-agent` is running with the role prompt, and runs complete or fail with receipts — not when the scheduler only created an `executing` row.

They are not yet: a full autonomous company that manages its own PRs, branches, or SA-2 coaching — that is the next execution slices on `dev`.

DashPro self-hosted CI — plain English

The one-sentence version

GitHub Actions for DashPro's heavy Android / Quality pipeline does **not** rent GitHub's cloud machines by default. It asks **your PC** (a registered "runner" named like `edudashpro-dashpro-2`) to do the work. If that runner is offline or busy, the Actions UI shows **Queued**.

How it works (kitchen analogy)

1. You push a branch or open a PR → GitHub posts a job on the Actions board.
2. The workflow says: "Only a machine tagged `self-hosted` + `dashpro` may take this."
3. The runner service on this host (`actions-runner-dashpro.service`) is the waiter that claims those jobs and runs the scripts (npm, Jest, Android tools, ...).

4. When the waiter is asleep (GitHub status **offline**) or already serving another table (**busy**), new jobs wait in line — that is the yellow **Queued** state.
5. Axon company agents (Cass / Soren / Dana) are **not** the GitHub runner. They are separate Cursor workers. The Actions “Agents” tab on github.com is also unrelated GitHub product UI.

Will this drain our RAM?

It can — and it has before. Quality Gates previously peaked near ~18G RAM / heavy swap and froze the desktop. That is why the runner unit is capped:

LIMIT	VALUE	MEANING
MemoryHigh	8G	Soft pressure — kernel starts reclaiming
MemoryMax	10G	Hard kill ceiling for the runner cgroup
MemorySwapMax	4G	Cap on swap the runner may use

Idle listener is usually well under **1G**. A Quality Gates / Unit Tests / Android job can climb toward the cap. **Do not** stack a full Android CI graph with Cursor open **and** a local Expo/EAS OTA export (OTA scripts also set `NODE_OPTIONS=--max-old-space-size=8192`). Prefer one heavy job at a time.

Check live pressure:

```
systemctl --user status actions-runner-dashpro.service
free -h
gh api repos/axon-control-ops/dashpro/actions/runners --jq '.runners[] | {name,status,busy}'
```

If systemd shows the service **running** but GitHub still says **offline**, the listener lost its session — restart the user unit, then confirm GitHub status flips to `online` :

```
systemctl --user restart actions-runner-dashpro.service
```

What “worker / run_...” means in the Actions header

That string is usually the **git branch name** from a company worker (`worker/run_<id>`), not proof that a runner is currently executing the job.

DashPro canary OTA vs EAS Build — why you may not see a build

Plain English

Canary OTA (`npm run ota:canary` in the DashPro project) publishes a **JavaScript/asset update** to the Expo `operator-canary` branch via `eas update` . That is **not** a native store binary rebuild.

YOU MIGHT LOOK FOR	WHAT CANARY OTA ACTUALLY IS
EAS Build (new APK/AAB/IPA)	No — not this path
EAS Update / Updates feed for branch <code>operator-canary</code>	Yes — after local export finishes
GitHub Actions “Android CI/CD Pipeline”	Separate — PR/push CI on the self-hosted runner
Team dock “Watch CI” / draft PR chips	GitHub delivery receipts — not Expo Update URLs

Why Dana’s thread looks busy but Expo still shows no build

Typical sequence for `ota:canary`:

1. Release guard checks the branch/gates.
2. `eas update ...` starts → **local** `expo export` (iOS + Android) with a large Node heap — this can take many minutes and uses ~1G+ RSS on the host.
3. Only after export succeeds does the update upload to Expo’s cloud.
4. Until step 3, the Expo dashboard has **no new Update** (and never an EAS Build for this command).

Watch progress on:

- Dana’s **IDE thread / terminal** (the live `expo export` / `eas update` lines)
- Expo dashboard → **Updates** (branch `operator-canary`), not Builds
- Host processes: `pgrep -af 'eas update|expo export'`

Axon Mission Control does **not** currently mirror Expo Update URLs the way it mirrors GitHub PR/CI links on the Team dock.

Quick reference commands

```
## CI locally
npm run verify:contracts
npm run verify:console-web

## Open PR
gh pr create --base dev --head feat/my-slice --title "..." --body "..."

## PR checks
gh pr checks

## DashPro self-hosted runner (this PC)
systemctl --user status actions-runner-dashpro.service
gh api repos/axon-control-ops/dashpro/actions/runners --jq '.runners[] | {name,status,busy}'

## Worker scheduler tests
PYTHONPATH=services/control-plane python3 -B -m unittest
tests.test_workspace_agent_scheduler -q

## Stale reconcile (when loaded on CP)
curl -X POST http://127.0.0.1:8787/api/runs/reconcile-stale
```

PDF: This chapter is bundled into `~/Desktop/Axon-X-How-To-Handbook.pdf` when you run `./scripts/docs/build-howto-handbook-pdf.sh` (rebuild after handbook edits).

40 Gate 9 — CI remediation (unaware-operator loop)

Updated: 2026-07-29

When **Axon-X Fast Gate** goes red, control-plane can ingest the GitHub `workflow_run` event, raise an inbox signal, lease a repair task to Rowan (watcher), one-shot dispatch a worker, and report the outcome — without you watching the Actions tab.

What is wired (v1)

STEP	BEHAVIOR
Ingest	<code>POST /api/webhooks/github/workflow-run</code> (HMAC <code>X-Hub-Signature-256</code>)
Config	<code>config/ci-remediation.json</code> — Axon-X Fast Gate and all six active DashPro workflows enabled
Dedup	One open repair per <code>(repo, workflow, branch, head_sha)</code>
Signal	Inbox item source <code>ci_remediation</code> merged in control-plane inbox projection
Lease	Task ledger goal starts with <code>CI repair:</code> · owner_role <code>watcher</code>
Dispatch	Background <code>dispatch_continuous_worker_run</code> when <code>dispatch_on_ingest</code> is true
Report	Worker posts <code>POST /api/ci-remediation/report-outcome</code> → signal + spoken line

Still human-gated: merge to `dev` / `master`, force-push, secrets.

GitHub webhook setup

1. Set `AXON_WATCH_GITHUB_WEBHOOK_SECRET` (or `GITHUB_WEBHOOK_SECRET`) on the control-plane host.
2. In GitHub → repo **Settings** → **Webhooks**: - Payload URL: `https://<control-plane>/api/webhooks/github/workflow-run` - Content type: `application/json` - Secret: same as the env var - Events: **Workflow runs**
3. Confirm delivery with a red Fast Gate push; inbox should show `Axon-X Fast Gate failed on <branch>`.

Local poll fallback (when public tunnel/DNS is down)

GitHub webhook deliveries fail with `connection_error` if `axon.edudashpro.org.za` does not resolve or the tunnel is degraded. Config alone cannot auto-fix CI in that state — Rowan never sees the event.

Use the local poller (HMAC to control-plane on localhost):

```
CONTROL_PLANE_URL=http://127.0.0.1:8787 \  
./scripts/ops/poll-fast-gate-remediation.sh
```

Always-on timer (user systemd):

```
mkdir -p ~/.config/systemd/user  
cp scripts/ops/systemd/axon-fast-gate-remediation-poll.{service,timer} \  
~/.config/systemd/user/  
systemctl --user daemon-reload  
systemctl --user enable --now axon-fast-gate-remediation-poll.timer
```

Public webhook remains preferred; the timer covers tunnel/DNS outages. The poller forwards both completed outcomes:

- **failure** → dedupe, Attention signal, leased **CI repair:** task, Rowan dispatch
- **success** → delivery update and stale failure cleanup; no repair is dispatched

Prove the local repair path is ready

```
systemctl --user is-enabled axon-fast-gate-remediation-poll.timer
systemctl --user is-active axon-fast-gate-remediation-poll.timer
systemctl --user show axon-fast-gate-remediation-poll.service -p Result
gh auth status
cursor agent status
curl -fsS http://127.0.0.1:8787/api/worker-scheduler | \
jq '{dispatch_enabled, env_allowed, effective_enabled}'
```

Required result: timer **enabled** + **active**, last service result **success**, authenticated GitHub and Cursor CLIs, and **dispatch_enabled: true**. **effective_enabled** controls ordinary scheduled shifts; webhook/poller repair dispatch uses the narrower **dispatch_enabled** gate.

Run the poller once to verify GitHub → signed control-plane ingest without waiting for the timer:

```
./scripts/ops/poll-fast-gate-remediation.sh
```

On a green head this is safe: it updates delivery/stale state and does not create a repair task.

Operator surfaces

- **Attention / inbox** — CI failure and repair/blocked updates
- **Task board** — leased **CI repair:** task for Rowan (escalate role Quinn)
- **Voice** — report-outcome returns a **spoken** line; ask-shaped dig-in on budget exhaustion
- **Stale clear** — when Fast Gate goes green again (webhook/poller) or on the worker scheduler tick, Axon-X confirms and resolves drill / superseded failure alerts so VAXON stops re-speaking them. Say “**clear stale alerts**” to VAXON to run the same confirmation sweep on demand.

What counts as confirmed-stale

SIGNAL	CLEARED WHEN
drill/... Fast Gate failure	Always (drill / useless after the exercise)
Real branch failure	Latest Fast Gate on that branch is success (via gh or green webhook)
Operator ack	Attention clear / acknowledge resolves Gate 9 CI store rows (not watch-only)

Stale CI-infra hint (branch predates a trunk fix)

Incident: `worker/run_96147146e311` (dashpro PR #39) burned its whole `attempt_budget` on Database SQL Linting without ever touching the real cause: its own `.github/workflows/db-lint.yml` still used `actions/setup-python@v5`, which cannot resolve a Python build for the `dashpro` self-hosted runner's OS and fails before any SQL is even linted. `development` had already fixed this (system-Python venv instead of `setup-python`) — the branch was just never rebased onto it. No content edit on the branch could have made this check pass.

`services/control-plane/app/ci_remediation/staleness.py` now checks for exactly this shape before a repair goal is built: if the failing workflow is currently **green on the PR's base branch**, and the base branch has commits touching CI-infra paths (`.github/workflows/**`, `.sqlfluff`, `scripts/check-migration-filenames.sh`, `scripts/workflow/**`, ...) that the failing head is missing, `build_repair_goal()` appends a hint telling the worker to merge the base branch in (fast-forward-safe merge commit, never force-push) **before** editing SQL/app content. The check is fail-open — any `gh` error, missing PR, or nothing-stale result silently omits the hint, it never blocks goal-building.

This is a hint, not a gate: the worker still decides what to do. It exists so the worker doesn't have to rediscover "merge trunk first" by trial and error across `attempt_budget` tries, the way `worker_96147146e311` did.

Related, separate fix landed from the same investigation: `Database SQL Linting` itself was failing on **every** PR touching `supabase/migrations/**` regardless of content, because `PG01` (CREATE INDEX CONCURRENTLY) flagged pre-existing indexes and the filename checker had no grandfathering for legacy names. See dashpro PR #66 — `.sqlfluff` now excludes `PG01` and `scripts/check-migration-filenames.sh` grandfathers a small, explicit `scripts/migration-filename-legacy-allowlist.txt`. A worker hitting a systemic, content-independent lint failure like this should fix the trunk config in its own small PR rather than trying to rewrite unrelated legacy migrations in its own branch.

Enabled workspaces

- Axon-X: `Axon-X Fast Gate`
- DashPro: `CI`, `Android CI/CD Pipeline`, `Database SQL Linting`, `Docs Policy Enforcement`, `Security Scan`, and `Voice Benchmark Nightly`

For another workspace, add an enabled binding in `config/ci-remediation.json` with the correct repository and workflow names. No new core code is required.

Local verify

```
./scripts/dev/python.sh -m unittest tests.test_ci_remediation
tests.test_ci_remediation_staleness -v
```

Manual red-path drill: use a throwaway branch and a deliberate Fast Gate ratchet overshoot, confirm webhook/poller → signal → leased task → Rowan repair draft PR → green head, then append evidence under `docs/ops/agent-reports/AUTONOMY-EVIDENCE-LOG.md`. Never drill on `dev`.

Autonomy gates & service identity (operator directives)

Do this in order. Do not skip ahead.

This page is the short, directive companion to [docs/AXON-X-AUTONOMY-MASTER-PLAN.md](#) . It tells you **what to turn on**, **what to leave off**, and **how to prove it**.

Last updated: 2026-07-25

Are we on auto-loop? No. See [auto-loop-and-credits.md](#) for the operator brief and Cursor credit budgets. Plumbing for Gates 6/9 and task-scope guardrails is in; unattended multi-project overnight loops are not.

1. Current gate status (what “done” means)

GATE	STATUS	YOU MUST
0–1	Closed	Keep a known-good commit; do not mix unrelated dirty trees into autonomy PRs
2 thin + residuals	Closed	Use <code>local_token</code> on any non-loopback surface; set operator + internal tokens
3	Closed	Continuous workers use disposable <code>worker/<run_id></code> worktrees only
4	Closed	Continuous workers only run from a leased task ; scheduler stays off until you intentionally enable it
5	Closed	Use Lead plan/fan-out/replan APIs; keep scheduler off unless you are watching a drill
6	Closed (thin)	Acceptance evidence / verifier must pass before completion/publish; task <code>allowed_paths</code> intersect contract scope
9	Proven (Axon-X)	Fast Gate red → CI repair task → draft PR → re-watch; supersede old repair chains; no protected merge

Scheduler rule (non-negotiable for daily driving): keep continuous workers `effective_enabled: false` unless you are deliberately drilling with a bounded task board and you are watching the machine. Enabling the scheduler is **not** the same as claiming a closed auto-loop — RETRY SHIFT / Critical Review failures still need human attention.

Check:

```
curl -sS http://127.0.0.1:8787/api/worker-scheduler | jq '{scheduler_enabled, effective_enabled, executing_count}'
```

Expect `effective_enabled: false` on the daily-driver host.

2. Gate 4 — Task ledger (how to use it)

What it is

Every continuous worker shift must bind to one durable task with:

- goal + acceptance criteria
- owner role
- attempt budget
- lease holder + expiry
- terminal outcome (`completed` / `failed` / `cancelled`)

Mission Control shows the board (open / leased / done / failed).

Operator actions

1. **Create a task** (API or Mission Control task board): - `POST /api/workspaces/{workspace_id}/tasks`
2. **Leave the scheduler off** unless you mean to run continuous workers.
3. When a worker claims work, the control plane **leases** the task, creates a run with that `task_id`, and refuses dispatch without a leased task.
4. Long shifts **renew** the lease at dispatch start so mid-run expiry does not silently drop ownership.
5. Dependencies must be `completed` before a dependent task can lease.

Do / Don't

DO	DON'T
Create explicit tasks before enabling continuous workers	Let workers “pick something useful” without a task ID
Keep attempt budgets small (default 3)	Leave infinite retries
Cancel obsolete open tasks when goals change	Leave stale open tasks for the wrong role
Verify with <code>tests.test_gate4_task_ledger</code>	Treat Gate 4 as done if only the UI exists

Proof commands:

```
./scripts/dev/python.sh -m unittest tests.test_gate4_task_ledger -q  
curl -sS http://127.0.0.1:8787/api/workspaces/workspace_demo/tasks | jq .
```

3. Gate 5 — Lead planner (how to use it)

Operator actions

1. Preview a goal as an ordered DAG: `POST /api/workspaces/{workspace_id}/lead/plan`
2. Materialize tasks and dependency-ready specialist runs: `POST /api/workspaces/{workspace_id}/lead/fan-out`
3. When the goal changes, use: `POST /api/workspaces/{workspace_id}/lead/replan` — do not manually leave obsolete tasks open.
4. After all specialist tasks are terminal, call: `POST /api/lead/plans/{plan_id}/synthesize`

Safety rules

- Lead assigns only company-roster specialist roles.
- Exact and parent/child path overlaps cannot lease concurrently.
- Explicit replans stop/cancel obsolete active work and persist receipts.
- “Check with all teammates” creates one leased task/run per specialist; it never picks one winner. It does not automatically open IDE tabs.
- Keep the continuous scheduler **off**. Gate 5 creates ready runs but does not authorize unattended verification.

Proof:

```
./scripts/dev/python.sh -m unittest \
  tests.test_lead_task_plan \
  tests.test_lead_fan_out \
  tests.test_lead_replan -q
```

4. Watch service identity (token + proxy mTLS)

Control-plane talks to watch on `/internal/watch/*`. That path must never be anonymous on a reachable host.

Minimum (required on remote)

Set the **same** shared token on control-plane and watch:

```
## ~/.config/axon-watch/deployment.env
AXON_WATCH_INTERNAL_SERVICE_TOKEN=<long random secret>
AXON_WATCH_REMOTELY_REACHABLE=1 # or a non-loopback AXON_WATCH_PUBLIC_BASE_URL
```

Do: generate with `openssl rand -hex 24`

Don't: leave the token empty on a tunnel / public URL

Don't: expose watch port `:8788` on the public internet

When remotely reachable, watch **refuses** mutating internal routes if the token is missing (HTTP 503) or wrong (HTTP 401).

Operator recovery (Mission Control **REQUIRED CONNECTOR DOWN** / vault unlock 503):

```
axonfixconnectors --ensure-internal-token --restart
## unlock Vault in the UI (Remember me), then:
axonfixconnectors
```

Trusted always-on host — re-enable vault auto-unlock without disabling remote auth:

```
## ~/.config/axon-watch/deployment.env
AXON_WATCH_ALLOW_VAULT_AUTO_UNLOCK=1
axonrestart
## /vault → Enable auto-unlock (keyfile already present will auto-unlock on next start)
```

Finish Cloudflare ingress cutover to `:4173` (ends soft-cutover `:7734` proxy):

```
## Cloudflare API token with Account → Cloudflare Tunnel → Edit
echo 'CF_API_TOKEN=...' >> ~/.config/axon-watch/deployment.env
./scripts/ops/set-tunnel-ingress-4173.sh
```

Full playbook: [docs/HOW-TO-HANDBOOK.md](#) .

Proxy mTLS capability (deployment proof required)

Axon-X supports proxy-verified client certificates **plus** the shared token. This is not end-to-end proof by itself: the proxy must strip incoming verification headers, verify the certificate, and be the only network path to watch.

1. Mint certs once:

```
./scripts/ops/mint-watch-mtls.sh
## writes ~/.config/axon-watch/mtls/{ca,client,server}.{crt,key}
```

2. Put this in `deployment.env` (both services):

```
AXON_WATCH_MTLS_REQUIRED=1
AXON_WATCH_MTLS_CA_FILE=$HOME/.config/axon-watch/mtls/ca.crt
AXON_WATCH_MTLS_CLIENT_CERT=$HOME/.config/axon-watch/mtls/client.crt
AXON_WATCH_MTLS_CLIENT_KEY=$HOME/.config/axon-watch/mtls/client.key
AXON_WATCH_MTLS_ALLOWED_CN=axon-control-plane
```

3. Configure the reverse proxy in front of watch to verify the client cert and forward:

- `X-SSL-Client-Verify: SUCCESS`
- `X-SSL-Client-S-DN: CN=axon-control-plane`

4. Restart: `axonrestart`

Do: keep token **and** mTLS together

Don't: enable `AXON_WATCH_MTLS_REQUIRED=1` without proxy verify headers (every mutating watch call will 401)

Don't: commit private keys into git

Proof:

```
./scripts/dev/python.sh -m unittest
tests.test_gate2_auth_containment.Gate2WatchInternalTokenTests -q
```

5. Control-plane operator auth (reminder)

On any remotely reachable console:

```
AXON_WATCH_AUTH_MODE=local_token
AXON_WATCH_OPERATOR_TOKEN=<strong secret>
AXON_WATCH_AUTH_ALLOW_LOOPBACK=0
```

Remote surfaces **force** `local_token` even if env still says `placeholder`. Full Access / exact-effect on remote also require step-up header `X-Axon-Step-Up` (`full-access` or `exact-effect`).

6. Verify before you claim a gate closed

```
npm run verify:contracts
npm run verify:console-web
## then push and confirm Axon-X Fast Gate is green on GitHub
```

Evidence lives under `docs/ops/agent-reports/` and the roll-up log `docs/ops/agent-reports/AUTONOMY-EVIDENCE-LOG.md`.

7. What to build next

1. Drive Critical Review + `Confidence: N/10` failure rate near zero (today's biggest auto-loop blocker in Mission Control)
2. Keep scheduler off by default; enable only for watched single-workspace drills
3. Harden multi-workspace work sources (CI repair + file-size patrol) under credit/RAM caps
4. Do not expand mobile mutation until this page's remote auth + mTLS steps are live on that host
5. Re-read `auto-loop-and-credits.md` before buying Ultra / enabling 3+ projects

49 Recent operator features (Gate 3–5 + console UX)

Updated: 2026-07-27

Audience: operators and agents who need plain-language “what changed and how to use it.”

Companion to [docs/HOW-TO-HANDBOOK.md](#) and [autonomy-gates-and-service-identity.md](#).

Mission Control second brain (2026-07-27): Transmission card, open vs focus workspace, fleet busy presence, and **REPORT** hotword — see [mission-control-transmission-and-report.md](#).

1. Mission Control task board (Gate 4)

Where it is (easy to miss)

Mission Control is the Operator **center panel** — but only in **Grid** view.

Your Brain Graph starfield (VAXON CORE / DashPro orbs) is Operator too; it **replaces** the Mission Control title, Fleet health, and Task board until you leave Graph.

YOU ARE HERE	HOW TO OPEN TASK BOARD
Top nav OPERATOR + Brain Graph	Click Mission Control in the bottom status chips, or Mission Control · Task board in the right Intelligence panel
Operator Grid already	Scroll under Fleet health → section Task board
Top nav IDE	Click OPERATOR first, then use Grid / Mission Control as above

On Grid you should see the heading **Mission Control**, then **GRID | BRAIN**, then **Fleet health**, then **Task board**.

What the task board is

A durable **task ledger**. Continuous workers may only start when they **lease** an open task. Chat in the IDE is separate (see “Talk vs do work” below).

How to use it

1. Top nav → **OPERATOR**.
2. Leave Brain Graph: **Mission Control** chip (or Intelligence → **Mission Control · Task board**).
3. Under **Fleet health**, use **Task board**.
4. Create a goal + owner role (e.g. [integrations](#)).

5. Leave the worker scheduler **off** unless you intentionally want continuous shifts.

Talk vs do work (IDE teammate modes)

In the IDE agent dock, open the mode chip (bottom-right, often **Agent FULL ACCESS**):

MODE	USE WHEN
Ask	Simple conversation — “thanks”, questions, status — no edits/tools
Plan	Map steps before changing code
Agent	Do the task — tools, edits, approvals
Debug	Reproduce and fix with evidence
Task board / Lead fan-out	Durable assigned work across specialists (Gate 4–5)

Examples

- “Thank you Dana!” / “I am thinking...” → stay on **Ask** (conversation).
- “Ok proceed with that.” / “Fix the payments race” → **Agent** (task).
- “Check with all sub-agents...” → Lead fan-out (N tasks/runs), not one specialty winner.

Example — seed work for Soren

```
curl -sS -X POST http://127.0.0.1:8787/api/workspaces/workspace_dashpro/tasks \
-H 'Content-Type: application/json' \
-d '{
  "goal": "Verify quality-gate heap calc does not OOM CI",
  "acceptance_criteria": "Targeted node-heap tests pass with receipts",
  "owner_role": "integrations",
  "attempt_budget": 2
}' | jq '{task_id, status, owner_role}'
```

Expect `"status": "open"`. When the scheduler is enabled, that role claims/leases it before `create_run`.

Example — list open tasks

```
curl -sS 'http://127.0.0.1:8787/api/workspaces/workspace_dashpro/tasks?status=open' \
| jq '.items[] | {task_id, goal, owner_role}'
```

2. Concurrent IDE conversation tabs

What changed

Switching conversation tabs **no longer stops** another teammate’s live stream. Each IDE thread keeps its own SSE session.

Example

- 1. Start an agent on **Soren**.
- 2. Open **Marco** and start (or watch) another run.
- 3. Switch back to Soren — his stream continues; the busy glow stays on both tabs.

Stop still applies to the **focused** thread only.

3. Brain Graph workspace labels + colors

What changed

Every **named workspace** orb shows a label (DashPro included). Orbs get a **muted stable tint** from `workspace_id` so they are easier to spot. Health tones (attention / critical) still override.

Example

Operator → Brain Graph → select DashPro in the left rail → the **DASHPRO** chip appears on its orb. Unlabeled spheres are usually signals/connectors/runs, not workspaces.

4. Lead planner + fan-out (Gate 5)

What it is

Dana (Lead) turns one goal + company roster into ordered plan items, **persists** them as leased `workspace_tasks`, and for fan-out opens **one run per ready specialist** (not a single specialty-route winner).

PIECE	MODULE / API
Pure plan	<code>build_lead_task_plan</code>
Persist	<code>persist_lead_task_plan</code> → Gate 4 ledger
Materialize	<code>materialize_lead_fan_out</code> → tasks + ready runs
HTTP	<code>POST /api/workspaces/{id}/lead/plan</code> · <code>.../lead/fan-out</code>

Continuous **dispatch** (Lane B) is still separate — fan-out creates leased tasks and **queued** runs with `lead_fan_out_assigned` receipts (not fake-executing). Assignment is posted into each specialist IDE thread. The scheduler (when on) promotes queued fan-out runs into Lane B; keep it off until Gate 6 is solid.

Examples

Fan-out (all specialists in parallel; path conflicts serialize / defer):

Goal: "Check with all sub-agents whether DashPro quality-gate heap calc is safe"
→ tasks for watcher, frontend, backend, integrations
→ N ready runs with receipts (deferred if exclusive_paths overlap)

```
curl -sS -X POST "http://127.0.0.1:8787/api/workspaces/workspace_axon_watch/lead/fan-out" \
-H 'content-type: application/json' \
-d '{"goal":"Check with all sub-agents whether Gate 5 fan-out is wired"}' \
| jq '{mode, run_count:(.runs|length), deferred:(.deferred|length), receipt}'
```

Sequential (then-chains → dependency task ids):

Goal: "Fix the API quality-gate heap calc then update the Expo confirmation screen"
→ backend task, then frontend task that depends on the backend task_id

Specialty route stays single-winner for normal prompts; fan-out phrases return `reason=lead_fan_out` instead of picking one teammate.

Proof:

```
./scripts/dev/python.sh -m unittest tests.test_lead_task_plan tests.test_lead_fan_out -q
```

4b. Continuous Lead takeover (specialist → Dana)

When a specialist shift ends (continuous worker **or** IDE Lane B), Dana now posts a **Lead takeover** rollup in the Lead IDE tab even if the shift was not linked to a formal Lead fan-out plan. Missing `task_id` no longer drops the handoff.

PIECE	LOCATION
Takeover report	<code>lead_takeover.post_lead_takeover_report</code>
Notify hook	<code>notify_lead_after_worker_task</code> (+ Lane B finalize)
Follow-up	Open <code>Lead follow-up after ...</code> task for Decide / next assign

Full multi-specialist synthesis still waits until every plan task is terminal; single specialist completions get an immediate Dana report so the loop stays continuous.

Why

`feat/*` and PRs to `dev` must keep **Axon-X Fast Gate** green. A red push blocks merge and often means file-size ratchets or unit tests.

Operator / agent command

```
## From repo root – polls the latest Fast Gate run for this branch
./scripts/ops/watch-fast-gate.sh
```

Or manually:

```
gh run list --branch "$(git branch --show-current)" --workflow "Axon-X Fast Gate" --limit 3
gh run watch # follows the newest in-progress run
gh run view --log-failed # when red
```

Who owns CI for which company

WORKSPACE	WHO WATCHES CI
Axon-X (<code>workspace_axon_watch</code>)	Rowan (watcher) + Mira (Lead triage) — Fast Gate for this repo
DashPro	Cass (watcher) + Soren (integrations / Actions)
Child workspaces	Watcher role for that company; escalate to Lead

VAXON / the Axon-X watcher should treat **this repository’s Fast Gate** as part of runtime health after any push to `origin`.

Gate 9 — unaware-operator remediation (Axon-X)

When Fast Gate fails, control-plane can ingest GitHub `workflow_run` events (HMAC webhook), raise an inbox signal, lease a repair task to **Rowan**, and report the outcome. See `ci-remediation-gate9.md`. Fallback poller: `./scripts/ops/poll-fast-gate-remediation.sh`.

5b. Lead team check-in scheduler (clocked)

Company Leads run a **cooldown-gated check-in** on the continuous worker clock (same tick family as file-size patrol — axon-local-style interval work).

PIECE	LOCATION
Work source	<code>lead_team_checkin</code> in <code>config/autonomy-work-sources.json</code>
Logic	<code>services/control-plane/app/workspace_agents/lead_team_checkin.py</code>
Hierarchy + close-out	<code>company-hierarchy-and-lead-checkin.md</code>

Failed shifts and degraded monitors/connectors become `Lead assigned:` tasks for the matching specialist. Usage/auth/restart failures escalate to the Lead IDE only (no bogus repair task).

External HTTP health probes: `http_health` checks in `config/axon-x-monitor-slice.json` and `config/dashpro-monitor-slice.json`.

6. Workspace company parity (vs DashPro)

DashPro is not a special UI — every bound workspace uses the same Team panel, Mission Control, and Task board. What differed was **staffing**:

WORKSPACE	COMPANY
DashPro / Axon-X / School	Named companies in <code>config/workspace-agents.json</code>
TPS / Young Eagles / Axon Local	Named companies (same shape as DashPro)
Other bound projects	Template staffing with workspace-stable names + voices (not Mira clones)

After switching to **TPS** in IDE → Team, you should see **Noor / Blair / Vera / Hugo / Tess**, not Mira/Rowan. Mission Control → Grid always keeps the **selected** workspace on the fleet board even when many projects exist.

7. What is still off on purpose

CONTROL	STATE	WHY
Continuous worker scheduler	Off by default	Unattended auto-loop not claimed yet — see auto-loop-and-credits
Live-checkout continuous edits	Forbidden	Gate 3 isolation — workers use disposable worktrees
Gate 5 full fan-out dispatch	Partial	Persist + ready runs land; Lane B auto-dispatch still often needs Retry / scheduler
Overnight multi-project loop	Off	Critical Review / dig-in / signal review still operator-gated

Check scheduler:

```
curl -sS http://127.0.0.1:8787/api/worker-scheduler \
  | jq '{scheduler_enabled, effective_enabled, executing_count}'
```

Auto-loop status & credits budget (operator brief)

Read this first when someone asks: “Are we autonomous yet?”

Last updated: 2026-07-25

1. Short answer

No — Axon-X is not yet on a closed unattended auto-loop.

What you have today is a **strong supervised company**: named employees, leased tasks, worktree isolation, Gate 6 acceptance, draft-PR delivery, Gate 9 CI repair, and (new) proactive file-size patrol *plumbing*. What you do **not** have is “Dana hands off, everyone runs overnight, and Mission Control only shows green in the morning” without operator nudges.

If Mission Control still shows any of these, you are in **supervised mode**:

UI SIGNAL	MEANING
RETRY SHIFT	A worker run failed; a human (or Lead) must restart it
Shall I dig in? / Yes–Not now	VAXON needs operator confirmation
Last shift failed: Critical Review Clause missing...	Shift ended without a valid Confidence: N/10 receipt
0 active runs + many open signals	Work is queued for review, not continuously draining
Typing “handoff the tasks” into Lead chat	Fan-out / lease pickup is not fully self-driving yet

That matches the DashPro screen on 2026-07-25: Marco failed Critical Review, Expo called out as down, Lead recommended next steps, and the operator had to type the handoff.

2. What is automatic today

LOOP PIECE	STATUS	NOTES
Create / lease durable tasks (Gate 4)	Ready	Continuous workers require a leased task
Lead plan → fan-out → specialist tasks (Gate 5)	Ready	Persist + ready runs; still often needs scheduler or Retry
Disposable worker worktrees (Gate 3)	Ready	Live checkout edits forbidden for continuous workers
Gate 6 verifier / acceptance evidence	Ready	Failed checks block completion/publish path
Draft PR delivery (no force- push)	Ready	Company pattern for publish
Gate 9 CI remediation (webhook → repair task)	Proven for Axon-X	Reactive self-heal when Fast Gate goes red
Per-task <code>allowed_paths</code> guardrails	Ready (new)	Hard scope for writes/publish + prompt anchors
File-size patrol work source	Plumbing ready	Scheduler can enqueue bounded hygiene tasks when enabled
Mission Control AUTONOMOUS attend	Implemented; live drill pending	Orb toggle → <code>autonomy_mode=full</code> ; safe findings dispatch; guarded findings require an exact approve/reject decision
Continuous worker scheduler	Off by default	Intentional — turn on via AUTONOMOUS ON / Settings Full
Unattended overnight multi- project loop	Not ready	Still needs human dig-in for Decide gates
Staging / production deploy autonomy	Not ready	Remains human-gated

Mission Control AUTONOMOUS button (orb)

CONTROL	MEANING
AUTONOMOUS ON	After confirmation, sets <code>autonomy_mode=full</code> ; workers run only when the host brake and dispatch gates allow
AUTONOMOUS OFF	Calls the hard-kill path: demotes to <code>semi</code> , pauses new starts, and attempts to stop active worker shifts
Hard-kill	Same emergency path; partial stop failures are reported and must be reviewed in active runs
Needs you receipts	Critical / secrets / production / protected merge / spend stay pending until Approve exact task or Reject

Auto-allowed while ON: inspect, retry idempotent checks, route internal handoffs, create bounded specialist tasks, edit disposable worktrees, run tests.

Always ask before mutation: critical severity, secrets/credentials, destructive filesystem/git/database, production deploy/release, protected merge/push, external/public communication, permissions/policy changes, approval-gated runs, raising usage/spend caps. Approval creates one `risk=approved`, attempt-bounded task linked to the decision receipt; rejection creates no task.

Findings dedupe while an open/leased task exists, while a decision is pending, or for a 15-minute receipt cooldown. Routed handoffs reuse their existing target task; only a handoff missing target routing may create a new attend task.

Status feed:

```
curl -sS http://127.0.0.1:8787/api/operator/autonomy/status \
  | jq '{autonomy_mode, autonomous_enabled, effective_autonomy, scheduler, last_scan,
pending: (.pending_critical_decisions|length)}'
```

Check the scheduler before assuming anything is looping:

```
curl -sS http://127.0.0.1:8787/api/worker-scheduler \
  | jq '{scheduler_enabled, effective_enabled, executing_count}'
```

Expect `effective_enabled: false` on the daily-driver host unless you deliberately enabled it for a drill.

3. What “auto-loop” must mean before we claim it

A workspace is on a real auto-loop only when **all** of these are true without operator chat prompts:

1. Open work is created from signals / Lead / hygiene sources (not only typed handoffs).
2. The scheduler claims leased tasks for the right roles on its own.

- 3. Shifts reliably end with Critical Review + **Confidence: N/10** .
- 4. Failed shifts reopen or escalate with a clear receipt (not a silent zombie lease).
- 5. Passing work becomes a **draft PR**; CI red creates a **bounded repair** (Gate 9).
- 6. Out-of-scope drift is blocked (allowed paths + scope guard).
- 7. You can leave the machine for hours and return to draft PRs / green repair heads — not a wall of RETRY SHIFT and dig-in prompts.

Until then: call it **bounded supervised autonomy**, not auto-loop.

4. Credits & subscription budget (Cursor + friends)

Axon continuous workers burn **Cursor Agent / CLI usage** (subscription pool or API key). Pricing moves; treat numbers below as **2026 order-of-magnitude** planning, and verify in the Cursor dashboard.

Cursor plans (individual / small team)

PLAN	ROUGH MONTHLY PRICE	INCLUDED MODEL USAGE (GUIDE)	FIT FOR AXON-X
Pro	~\$20	~\$20 API-equivalent	Operator + occasional IDE agent; not multi-project continuous
Pro+	~\$60	~\$70	1 proving workspace, light scheduler drills, prefer Auto
Ultra	~\$200	~\$400	Serious continuous agents on 1–2 busy workspaces
Teams	~\$40/user	Per-seat allowance	Shared billing / analytics; still need Ultra-class usage for many agents

Sources to re-check: [Cursor pricing help](#), [Models & pricing](#).

Rule of thumb: prefer **Auto** / cheaper router models for patrol and routine shifts; save frontier models for Lead synthesis and hard repairs. Frontier / MAX-style sessions can empty even Ultra in days.

How much autonomy burns (planning ranges)

Assume one “shift” ≈ one continuous worker run with tools (often many model turns). Costs vary wildly by model and repo size.

OPERATING MODE	SHIFTS / DAY (ALL PROJECTS)	SUGGESTED CURSOR TIER	MONTHLY CURSOR BUDGET TO PLAN
Supervised operator (today's default)	2–8 manual / Retry	Pro or Pro+	\$20–\$100
Bounded auto-loop drill (1 workspace, scheduler on, watched)	10–25	Pro+ → Ultra	\$100–\$250
Multi-project autonomous (3–6 workspaces, CI repair + patrol)	30–80	Ultra + on-demand, or Teams + pooled	\$250–\$800+
Aggressive 24/7 multi-agent (not recommended yet)	100+	Ultra + hard spend caps	\$800–\$2,000+

These are **planning envelopes**, not invoices. Track real burn for one week of drills before buying more seats.

Other credits you will also need

SURFACE	WHY IT MATTERS	ROUGH PLANNING NOTE
GitHub Actions	Fast Gate / child CI on every draft PR	Private repos: minutes; keep repair PRs small
Supabase / cloud DBs	DashPro storage already ~60% of 1 GB in the UI snapshot	Budget storage + egress; don't let agents spam uploads
Sentry / PostHog	Watch signals drive “dig in”	Event volume rises with more projects
Azure TTS / voice	Talk / briefing	Small vs Cursor; still meter it
Host RAM / OOM	Control-plane can die (exit 137) under concurrent agents	Cap <code>max_active</code> (default 2); more projects ≠ more concurrent agents

Practical recommendation (multi-project)

1. Stay **supervised** until Critical Review failures and RETRY SHIFT noise are rare.
2. Run **one proving workspace** (DashPro or Axon-X) with scheduler on for short drills.
3. Budget **Ultra** (~\$200–400/mo usage) before enabling continuous workers on 3+ projects.
4. Set a Cursor **spend limit** so on-demand cannot surprise you.
5. Keep Axon `max_active_executing` at 2 (or 1) — credits and RAM both explode with concurrency.

5. Operator checklist (today)

1. Confirm control-plane health: `curl -sS http://127.0.0.1:8787/api/health`

2. Confirm scheduler off unless drilling: see §2 command.
3. Clear failed Critical Review shifts (Retry or cancel stale tasks).
4. Prefer explicit tasks / Lead fan-out over free-form “go ahead and handoff” chat.
5. After any push: `./scripts/ops/watch-fast-gate.sh`
6. Re-read this page after each autonomy gate change.

Related:

- [autonomy-gates-and-service-identity.md](#)
- [recent-operator-features.md](#)
- [docs/AXON-X-AUTONOMY-MASTER-PLAN.md](#)
- [docs/AXON-X-AUTONOMY-READINESS.md](#)

52 Company hierarchy, Lead check-in, and monitoring close-out

Audience: operator + company Leads (Mira, Dana, Lindi, ...)

Clock: continuous worker scheduler tick (`services/control-plane/app/workspace_agents/scheduler.py`)

Work source: `lead_team_checkin` in `config/autonomy-work-sources.json`

Confirmed hierarchy flow

```
You (Decide)
  ↑↓
VAXON (fleet brain / operator thread)
  ↑↓
Company Lead (Mira / Dana / Lindi / ...)
  ↑↓
Specialists (watcher · frontend · backend · integrations)
  ↑
Lead rollup → VAXON → You for Decide gates
```

LAYER	OWNS
You	Priorities, Decide approvals, production go/no-go
VAXON	Cross-company routing, spoken alerts, Lead charter intents
Lead	Team check-in, assign the right specialist, rollup to VAXON
Watcher	Signals, CI/Fast Gate, Sentry/Vercel health, CRC completion
Frontend	Console / Expo / UI surfaces
Backend	APIs, control plane, persistence
Integrations	Connectors, PostHog/Supabase/email wiring, cross-repo

Axon-X (`workspace_axon_watch`) coaches other Leads; child companies own their product trees. EduDash Pro marketing site (`edudashpro.org.za`) is **DashPro** (Dana / Cass / Priya / Soren) — not EDP Excellence.

Lead team check-in (clocked guardrails)

Same pattern as axon-local's interval scheduler: the continuous worker **tick** (~45s) runs explicit work sources. Lead check-in is cooldown-gated (default **900s / 15 min** per company) so it does not spam Lane B.

Each due company Lead pass:

1. Scans **failed specialist shifts** (roster last outcome).
2. Scans **watch monitors + required connectors** (third-party / public health).
3. **Assigns** the matching specialist via leased ledger tasks (`Lead assigned:`).
4. **Escalates only** (Lead IDE note, no code task) for usage limits, runtime auth, restart interrupts, or operator-stopped shifts.
5. Posts a short check-in note into the Lead IDE thread when there is work.

Assignment guardrails

SIGNAL	OWNER
Failed watcher / frontend / backend / integrations shift	Same role
Sentry / CI / GitHub / Vercel monitor	watcher
PostHog / Supabase / email / generic connector	integrations
Control-plane / API HTTP health	backend
Console / UI HTTP health	frontend
Usage limit / vault auth / restart interrupt	Escalate to Lead → VAXON (no auto repair task)

Continuous workers still pick up `Lead assigned:` / `VAXON attend:` tasks when the fleet scheduler is enabled; Leads stay `on_demand` and do not auto-start as workers.

VAXON attend loop (Full autonomy)

When Mission Control **AUTONOMOUS ON** (`autonomy_mode=full`), the scheduler also runs work source `autonomous_attention` :

1. Collects failed shifts, monitors/connectors, inbox warnings, and open handoffs.
2. Classifies each item with a fail-closed safety policy.
3. **Dispatches** auto-safe items as deduped specialist tasks (`VAXON attend:`).
4. Reuses an already-routed handoff target task instead of creating a duplicate.
5. **Escalates** critical/dangerous items into pending autonomy decisions.
6. **Approve exact task** creates one attempt-bounded `risk=approved` task; **Reject** resolves the decision without work.
7. High / unclassified task `risk` cannot be claimed by continuous workers.

Emergency stop: Mission Control **Hard-kill** (or **AUTONOMOUS OFF**) demotes Full → Semi, pauses new starts, and attempts to stop active worker shifts. Any partial stop failures remain visible for operator review.

Manual fan-out remains available:

```
curl -sS -X POST "http://127.0.0.1:8787/api/workspaces/workspace_dashpro/lead/fan-out" \
-H 'content-type: application/json' \
-d '{"goal":"Check with all sub-agents on current blockers"}'
```

External + third-party monitoring

SURFACE	CONFIG
DashPro Sentry / PostHog / Supabase + site/GitHub HTTP	<code>config/dashpro-monitor-slice.json</code>
Axon-X public origin / control plane / GitHub HTTP	<code>config/axon-x-monitor-slice.json</code>
Connector probes (required + optional third-party)	<code>config/watch-connectors.json</code>
New check type	<code>http_health</code> in <code>services/axon-watch/app/monitors/http_health.py</code>

Inbox still projects non-ok monitor statuses; Lead check-in turns those into role-owned tasks on the clock.

Continuous specialist → Lead → VAXON

When a specialist finishes an IDE or continuous job (even without a Lead plan):

1. **Specialist** ends with a short handoff: what changed, verified, Blockers / Lead next.
2. **Lead** gets an automatic takeover note in the Lead IDE tab + a follow-up task.
3. **VAXON** gets a short operator-thread flash so `REPORT` / update has live fleet memory.
4. Full Lead **plans** still synthesize into a richer VAXON handoff when all plan tasks finish.

```
Specialist finishes
→ Lead takeover (Dana tab)
→ VAXON flash (operator thread)
→ You ask REPORT → VAXON uses roster + those handoffs
```

Close-out checklist — monitoring mode with Axon-X in play

Use this when shifting from build/ship into supervised monitoring:

1. **Fast Gate green** on the active Axon-X branch (`./scripts/ops/watch-fast-gate.sh`).
2. **Public origin** healthy (`https://axon.edudashpro.org.za/api/health`).
3. **Worker scheduler** on only when you want continuous specialists + Lead check-in assignments to dispatch (Mission Control → fleet scheduler).
4. **CI remediation** webhook or poller live (`docs/how-to/ci-remediation-gate9.md`).
5. **Vault** holds Sentry/PostHog/Supabase tokens so DashPro monitors are not stuck in `skipped` .

6. Confirm company roster in `config/workspace-agents.json` matches the hierarchy above.

7. Leave Decide gates with you; VAXON + Leads escalate, they do not silently ship production.

Proof:

```
./scripts/dev/python.sh -m unittest \
  tests.test_lead_team_checkin \
  tests.test_http_health_monitor \
  tests.test_autonomy_guardrails \
  -q
```

Confidence: 8/10 — clocked Lead assignment + HTTP monitor hardening are wired and unit-tested; live multi-company assignment under load still depends on scheduler enablement and vault credentials.

53 Recovering a stuck agent and running School Operations

Phase 1

Updated: 2026-08-07

This chapter answers two practical questions:

1. What should happen when an Axon-X employee agent is stuck or failed?
2. How should the first Young Eagles homework workflow operate safely?

It is written for an operator. VAXON should explain the reason and recommend the next safe action; it should not turn an Ask-mode conversation into work or retry a failure blindly.

A stuck agent: read the reason before retrying

Open the agent's report and use the exact failure text to choose the next step. The label **FAILED** tells you the outcome, not the cause.

WHAT THE REPORT SAYS	WHAT IT MEANS	YOUR NEXT ACTION	WHAT VAXON SHOULD DO
<code>control-plane</code> restart, <code>cancelled</code> after restart, or a run vanished while the service restarted	The platform interrupted the shift; it is not evidence that the agent's work failed.	Wait for health to return, then check that the task is Open , not stuck Leased . Retry only after that.	Preserve the cancelled-run receipt; release the exact interrupted lease without using up an attempt; explain that the retry is safe.
<code>sandbox</code> mount, <code>approved</code> writable root, or read-only filesystem	The worker did not get the workspace/scope it needed.	Do not keep pressing Try again . Report the exact error or ask VAXON to diagnose it.	Classify it as a runtime/configuration fault; repair or route the task to an editing specialist with a narrow writable scope, then retry once.
<code>consultative-only</code> , no tool access, or no files changed	The task was sent to a conversational/observer context rather than an execution-capable specialist.	Keep Ask for advice; use Assign or an execution task for implementation.	Explain the distinction, preserve the advice as a plan, and route the bounded implementation to the right role.
Confidence: <code>N/10</code> missing, missing receipt, or Gate 6/critical review failure	The evidence is incomplete or the result is not trustworthy yet.	Review the named check and the changed files. Do not approve it just to clear the card.	State the missing proof, request a concrete correction, and retry only when the evidence requirement is satisfiable.
missing credential, approval, parent decision, or external service failure	An authorised human or external system must act.	Supply the missing approval/credential through the approved channel, or decide to stop the task.	Mark the task <code>awaiting_input</code> / <code>waiting_external</code> , name exactly what is needed, and avoid speculative retries.
a real test/assertion failure	The implementation needs repair.	Use Explain first if the failure is unclear; use Try again only after the agent identifies a changed, testable fix.	Summarise the failing check, propose a bounded fix and test, then hand it to the owning specialist.

Restart decision: do not restart by default

Run `axonhealth` first. If it reports healthy and the UI is merely showing a failed agent card, restarting the control plane is usually the wrong first action: it interrupts other in-flight work.

Restart only when one of these is true:

- health/API is unavailable or the shell is genuinely empty/hung;
- a deployed backend change explicitly needs a reload; or
- VAXON's diagnosis identifies a control-plane fault and there are no active shifts that must be allowed to finish.

On this host:

```
axonhealth      # inspect first
axonrestart     # controlled soft restart when the service is the problem
axonrevive      # only for a wedged/unavailable runtime
```

After a restart, wait for `axonhealth` to return healthy, refresh the browser, then verify the affected task's state before retrying it. A restart-interrupted role task should return to **Open** with its retry budget intact. If it remains **Leased**, stop retrying and report the task ID and run ID to VAXON: that is a runtime recovery defect, not operator error.

VAXON's recovery contract

For every failed or stuck employee, VAXON should provide a short report in this order:

1. **What happened** — the exact class of failure in plain language.
2. **What changed / did not change** — files, tests, and receipts; never claim evidence it did not obtain.
3. **Who acts next** — operator, named specialist, or external owner.
4. **One safe next action** — explain, repair, request input, or retry.
5. **Confidence: N/10** — only after the summary.

It should not create a loop of retries. Two genuine, corrected attempts are a reasonable escalation point; infrastructure interruptions do not count as a genuine agent attempt.

Young Eagles Phase 1: daily homework with human approval

Phase 1 is deliberately narrow: publish daily homework safely. It is not an authorisation for an AI to issue grades, send parent messages, or make child welfare decisions on its own.

Daily operating flow

1. The teacher prepares the room-specific homework in **Teacher Workspace** and selects **Submit homework**.
2. The post enters **pending approval**. Parents must not receive it at this stage.
3. The principal or authorised administrator checks age group, wording, date, learning activity, and any attached practice material.

4. Approval changes the post to **posted** for the relevant room/age group.
5. At close of day, the team records what was posted and prepares the next draft; missed posts go onto the next opening checklist.

The existing Young Eagles Command Centre implementation follows this shape: homework posts create approval entries, and an approved homework entry is published for its room. Its current test suite covers syntax, role navigation, login/RBAC, API behaviour, and the related operations flows. It does **not** prove that an AI-generated message is educationally appropriate; the approving human remains accountable.

What Imani may do in this phase

Imani can act as a school-operations coordinator:

- prepare drafts from the teacher's supplied lesson plan;
- remind staff about missing homework, approvals, and weekly updates;
- prepare a practice-test or report draft from approved source records;
- summarise trends for a teacher, principal, or parent meeting; and
- suggest aftercare activities from the approved programme.

Imani must not, without the named human approval path:

- publish homework or parent communications;
- assign final marks, diagnose a child, or make promotion/disciplinary decisions;
- expose another child's information in a report or message; or
- make safety, pickup, medical, or safeguarding decisions.

Ready-for-Phase-2 checklist

Do not automate grading, reports, parent updates, or aftercare decisions until the centre has named owners, approval checkpoints, record-retention rules, and a clear parent escalation route for each of them. Start with one supervised room and review a week's audit trail before extending the workflow.

54 Reliability and deliberate controls

Updated: 2026-07-29

This chapter records the operator-facing reliability contract for the local console, Command Theater, voice playback, and IDE AgentDock.

Control-plane restart contract

- `control-plane.service` uses `Restart=always` and a short stop timeout.
- Its relationship with `axon-watch.service` is soft: a watch-service flap must not stop the control-plane.
- `axonrestart` restarts the control-plane first and waits for `:8787` health before continuing through the remaining units.
- Use `axonrevive` when the process is wedged rather than merely restarting.

Vite control-plane recovery

Start the source console through the guarded wrapper:

```
./scripts/ops/run-5173.sh
```

The wrapper checks `http://127.0.0.1:8787/api/health`, starts `control-plane.service` when necessary, and waits before starting Vite. If the API restarts while Vite remains up, `/api/*` returns a temporary JSON `503` instead of exposing `ECONNREFUSED`, then recovers to `200`.

```
./scripts/verify/verify-vite-cp-restart.sh
```

If recovery fails:

```
systemctl --user status control-plane.service  
journalctl --user -u control-plane.service -n 80  
axonrevive
```

Deliberate STAND-UP

Command Theater must open only from the STAND-UP control or an explicit VAXON request. Briefing polls do not auto-start it. The browser smoke check asserts that the theater is closed before the operator clicks.

Speech onset and stage synchronization

Azure output includes backend lead silence, while the browser player waits for audio readiness and preroll. The report stage changes from the audible-start callback, not from fetch completion, so the visible speaker and first audible word stay aligned. Browser speech uses its own short start guard.

AgentDock operator-message actions

In IDE AgentDock, **Edit** and **Resend** reserve their layout space but remain visually hidden until:

- the operator message is hovered with a fine pointer, or
- keyboard focus enters either action.

Touch and coarse-pointer devices keep the controls visible because they have no reliable hover state. Reduced-motion users get the same visibility behavior without the transition.

Verify the structural and interaction contract:

```
./scripts/dev/python.sh scripts/verify/check_dock_behavior_contract.py  
npm run typecheck -w @axon-watch/console-web
```

VAXON Desktop

Packaged Linux desktop shell (Tauri 2) with frozen Watch + Control Plane sidecars on loopback `:8787`.

Verify

```
npm run verify:desktop  
npm run prove:desktop:clean-install  
npm run prove:desktop:voice
```

Evidence: [docs/VAXON_DESKTOP_VERIFY_EVIDENCE.md](#).

Install

```
npm run build:desktop:linux  
sudo apt-get install -y ./apps/console-desktop/src-  
tauri/target/release/bundle/deb/VAXON_*.deb
```

Requires Debian-family x86_64 + WebKitGTK 4.1; current sidecars need **glibc** ≥ 2.42 . Full first-run / tray / update notes: [docs/VAXON_DESKTOP_VERIFY_EVIDENCE.md](#) and [docs/DUAL_RUNTIME_STARTUP_CONTRACT.md](#).

55 Online research: SearXNG and legacy Google

Preferred → fallback order (matches the control-plane research path):

1. **SearXNG** when `AXON_WATCH_SEARXNG_URL` is set (recommended local path)
2. **Google Custom Search** (legacy) when `AXON_WATCH_GOOGLE_CSE_API_KEY` and `AXON_WATCH_GOOGLE_CSE_CX` are set (DashPro aliases `EXPO_PUBLIC_GOOGLE_CSE_*` still work)
3. **DuckDuckGo Instant** as the last resort (often sparse)

Preferred: local SearXNG

Start the local instance and point research at it:

```
./scripts/dev/run-searxng.sh
## then set AXON_WATCH_SEARXNG_URL=http://127.0.0.1:8080 in .env (or vault)
```

`config/searxng/settings.yml` must keep `json` under `search.formats` (SearXNG returns **403** for `format=json` when JSON is disabled). Do not rely on public instances that disable the JSON API. After a successful search, capability / receipts should show `provider: searxng`.

SearXNG Search API: https://docs.searxng.org/dev/search_api.html

Legacy: Google Custom Search 403

If SearXNG is not configured and live search reports **403** — **this project does not have access to Custom Search JSON API** (reason often `forbidden`):

1. Open [Google Cloud Console → APIs & Services → Library](#) for the project that owns that API key (DashPro's Firebase/Cloud project is often `edudashpro`).
2. Find **Custom Search API** (`customsearch.googleapis.com`) and click **Enable**.
3. In **APIs & Services → Enabled APIs / Dashboard**, confirm Custom Search API is listed as enabled for that same project (not a different Cloud project).
4. Confirm the Programmable Search Engine ID (`cx`) still matches the engine in the [Programmable Search control panel](#).
5. Retry a search, or prefer local SearXNG instead of unblocking Google.

If Enable was already clicked and the 403 persists, Google's docs state the Custom Search JSON API is **closed to new customers** (existing customers only, through 2027). In that case, enabling the library toggle cannot grant access; use SearXNG (`./scripts/dev/run-searxng.sh`) rather than creating a new Google search project for Axon-X.

Official Google overview (legacy path only): <https://developers.google.com/custom-search/v1/overview>